

HW3 — Fitting Models and Image Wrapping

1 RANSAC and Fitting Models

1.1 Task1: RANSAC Theory

1.1.1 Minimum # of points

To compute a putative model, the minimum number of 3D points needed to sample in an iteration is three.

1.1.2 Probability single iteration fails

The probability P that the data picked for the putative model in a single iteration fails, assuming an outlier ratio $e = 0.5$, is calculated as:

$$P = 1 - (1 - e)^s$$

Where $s = 3$ for the case of 3D planes. Plugging the values we get:

$$P = 1 - (0.5)^3 = 0.875$$

Thus, the probability of failure in a single iteration is 87.5

1.1.3 Minimum # of RANSAC trials

The minimum number of RANSAC trials n needed to achieve at least a 98% chance of success P , with an outlier ratio $e = 0.5$, is given by the formula:

$$1 - (1 - (1 - e)^s)^n \geq P$$

Rearranging for n , we get:

$$n \geq \frac{\log(1 - P)}{\log(1 - (1 - e)^s)}$$

Plugging in the values for $P = 0.98$ and $s = 3$, we find:

$$n \geq \frac{\log(1 - 0.98)}{\log(1 - (0.5)^3)}$$

Upon calculation, we find:

$$n \geq 30$$

Therefore, at least 30 trials are needed to have a 98% chance of success.

1.2 Task2: Fitting Linear Transformations

1.2.1 Degrees of freedom, Minimum # of correspondences

The matrix M representing a linear transformation in $\mathbb{R}^{2 \times 2}$ has four degrees of freedom since it can be parameterized as $M = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$. Therefore, to fully constrain or estimate M , we need a minimum of two 2D correspondences.

1.2.2 Form of A, m, and b

Given 2D correspondences $(x'_i, y'_i)^T \leftrightarrow (x_i, y_i)^T$, we formulate the fitting problem as a least-squares problem:

$$\arg \min_{m \in \mathbb{R}^4} \|Am - b\|^2$$

where $m = [a, b, c, d]^T$ contains all the parameters of M , A is a $2N \times 4$ matrix dependent on the points (x_i, y_i) , and b is a $2N \times 1$ vector containing the coordinates (x'_i, y'_i) . The matrices are defined as:

$$A = \begin{bmatrix} x_1 & y_1 & 0 & 0 \\ 0 & 0 & x_1 & y_1 \\ x_2 & y_2 & 0 & 0 \\ 0 & 0 & x_2 & y_2 \\ \vdots & \vdots & \vdots & \vdots \\ x_N & y_N & 0 & 0 \\ 0 & 0 & x_N & y_N \end{bmatrix}, \quad b = \begin{bmatrix} x'_1 \\ y'_1 \\ x'_2 \\ y'_2 \\ \vdots \\ x'_N \\ y'_N \end{bmatrix}$$

The solution m that minimizes the sum of squared differences is found by solving the normal equations:

$$m = (A^T A)^{-1} A^T b$$

This provides the parameters of the linear transformation that best fits the correspondences in the least-squares sense.

1.3 Task3: Fitting Affine Transformations

1.3.1 Code implementation and Report (S, t)

For points_case_1.npy:

$$S = \begin{bmatrix} 1.41444296 & -1.41424374 \\ -0.70762108 & -0.70690933 \end{bmatrix} \quad (1-1)$$

$$t = \begin{bmatrix} 0.09998617 \\ 0.20014656 \end{bmatrix} \quad (1-2)$$

```

1 def p3(filename: str):
2     # code for Task 3
3     # 1. load points X from task3/
4     X = np.load(filename)
5     N, D = X.shape
6
7     # 2. fit a transformation  $y=Sx+t$ 
8     X_train = X[:, :2]
9     y_train = X[:, 2:]
10
11     A = np.zeros((2*N, D+2))
12     A[:N, :2] = np.copy(X_train)
13     A[N:, 2:4] = np.copy(X_train)
14     A[:N, 4] = 1.0
15     A[N:, 5] = 1.0
16
17     b = np.zeros((2*N, 1))
18     b[:N, 0] = y_train[:, 0]
19     b[N:, 0] = y_train[:, 1]
20
21     Result = np.linalg.lstsq(A, b)
22     S = Result[0][:4].reshape(2, 2)
23     t = Result[0][4:]
24     # print(S, t)
25
26     # 3. transform the points
27     X_transformed = (S @ X_train.T + t).T
28
29     # 4. plot the original points and transformed points
30     plt.scatter(X[:, 0], X[:, 1], label='Original points', c='blue', s=1)
31     plt.scatter(X[:, 2], X[:, 3], label='Transformed GT', c='red', s=1)
32     plt.scatter(X_transformed[:, 0], X_transformed[:, 1], label='
Transformed points', c='green', s=1.5)
33     plt.legend()
34     plt.show()
35
36     return S, t

```

1.3.2 Plots




Figure 1: Plot for points_case_1.npy




Figure 2: Plot for points_case_2.npy

1.3.3 Discussion

Affine transformation describe the data transformation in points_cast_1.npy quite good, but performs bad for points_cast_2.npy. This is because points_cast_1.npy is affine transformation, which can be described with 6 degree of freedom. But points_cast_2.npy is projection transformation, or homography, this transformation need higher dimension to describe.

1.4 Task4: Fitting Homographies

1.4.1 Implementation for fit_homography() in homography.py

Submitted to Canvas.

```
1 """
2 Homography fitting functions
3 You should write these
4 """
5 import numpy as np
6 from common import homography_transform
7
8 def fit_homography(XY):
9     '''
10     Given a set of N correspondences XY of the form [x,y,x',y'],
11     fit a homography from [x,y,1] to [x',y',1].
12
13     Input - XY: an array with size(N,4), each row contains two
14             points in the form [x_i, y_i, x'_i, y'_i] (1,4)
15     Output -H: a (3,3) homography matrix that (if the correspondences can
16     be
17             described by a homography) satisfies [x',y',1]^T == H [x,y
18     ,1]^T
19
20     '''
21     N = XY.shape[0]
22     A = np.zeros((2*N, 9))
23     for i in range(N):
24         x, y, xp, yp = XY[i]
25         # A[2*i] = [-x, -y, -1, 0, 0, 0, x*xp, y*xp, xp]
26         # A[2*i+1] = [0, 0, 0, -x, -y, -1, x*yp, y*yp, yp]
27         A[2*i] = [0, 0, 0, -x, -y, -1, x*yp, y*yp, yp]
28         A[2*i+1] = [x, y, 1, 0, 0, 0, -x*xp, -y*xp, -xp]
29
30     # Perform Singular Value Decomposition (SVD)
31     U, S, Vt = np.linalg.svd(A)
32
33     # The solution is the last column of V (or the last row of V transpose
34     )
35     h = Vt[-1]
36     # Normalize h
37     h /= np.linalg.norm(h)
38     # Reshape h to get the homography matrix H
39     H = h.reshape(3, 3)
40     return H
```

1.4.2 Report H

H for case 1 and case 4:

$$H1 = \begin{bmatrix} 1.00555949e+00 & 1.61370672e-03 & -1.35143989e-01 \\ 2.56045861e-03 & 6.22536404e-01 & -7.35872070e-01 \\ 4.51704286e-05 & 3.59823762e-05 & 1.00000000e+00 \end{bmatrix} \quad (1-3)$$

$$H4 = \begin{bmatrix} 1.63877010e-14 & 1.00000000e+00 & 1.04256051e-13 \\ 1.00000000e+00 & 1.03129045e-16 & 6.32521301e-14 \\ 3.68894767e-17 & 9.67901093e-17 & 1.00000000e+00 \end{bmatrix} \quad (1-4)$$

1.4.3 Plots




Figure 3: Plot for points_case_5.npy

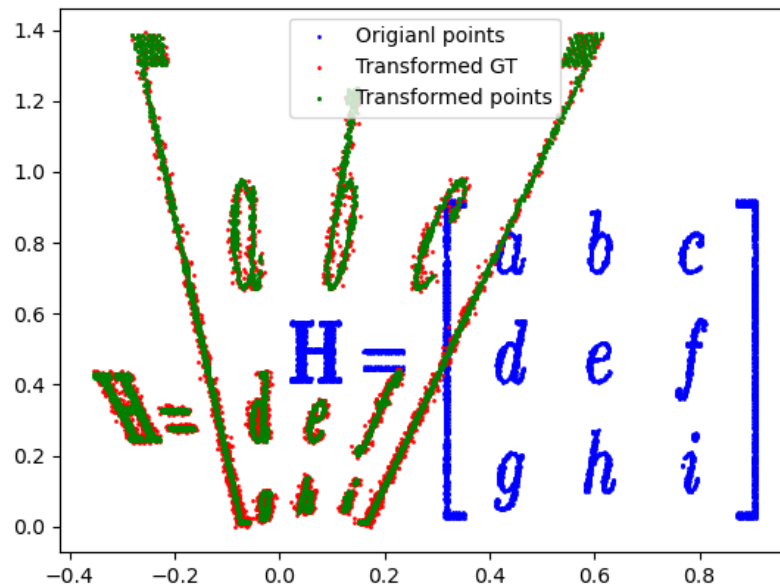


Figure 4: points_case_9.npy

2 Image Warping and Homographies

2.1 Task5: Synthetic Views - Name that Book!

2.1.1 Code implement of make_synthetic_view(sceneImage, corners, size) in task5.py

Code submitted to Canvas.

```

1  """
2  Task 5 Code
3  """
4  import numpy as np
5  from matplotlib import pyplot as plt
6  from common import save_img, read_img
7  from homography import fit_homography, homography_transform
8  import os
9  import cv2
10
11
12  def make_synthetic_view(img, corners, size):
13      """
14      Creates an image with a synthetic view of selected region in the image
15      from the front. The region is bounded by a quadrilateral denoted by
16      the
17      corners array. The size array defines the size of the final image.
18
19      Input - img: image file of shape (H,W,3)
20              corner: array containing corners of the book cover in

```

```

20         the order [top-left, top-right, bottom-right, bottom-left]
        (4,2)
21         size: array containing size of book cover in inches [height,
        width] (1,2)
22
23     Output - A fronto-parallel view of selected pixels (the book as if the
        cover is
24         parallel to the image plane), using 100 pixels per inch.
25     """
26     # The desired coordinates for the book corners
27     h, w = size
28     # Convert from inches to pixels: 1 inch is 100 pixels
29     h, w = h * 100, w * 100
30     dst_points = np.array([[0, 0], [w - 1, 0], [w - 1, h - 1], [0, h -
1]], dtype='float32')
31     XY = np.hstack((corners, dst_points))
32
33     # Compute the homography matrix
34     h_matrix = fit_homography(XY)
35
36     # Perform the warp perspective
37     warped_image = cv2.warpPerspective(img, h_matrix, (int(w), int(h)))
38     return warped_image
39
40 if __name__ == "__main__":
41     # Task 5
42
43     case_name = "threebody"
44
45     I = read_img(os.path.join("task5", case_name, "book.jpg"))
46     corners = np.load(os.path.join("task5", case_name, "corners.npy"))
47     size = np.load(os.path.join("task5", case_name, "size.npy"))
48     # import pdb; pdb.set_trace()
49
50     result = make_synthetic_view(I, corners, tuple(size[0]))
51     save_img(result, case_name+"_frontoparallel.jpg")
52
53
54     case_name = "palmer"
55
56     I = read_img(os.path.join("task5", case_name, "book.jpg"))
57     corners = np.load(os.path.join("task5", case_name, "corners.npy"))
58     size = np.load(os.path.join("task5", case_name, "size.npy"))
59     # import pdb; pdb.set_trace()
60
61     result = make_synthetic_view(I, corners, tuple(size[0]))
62     save_img(result, case_name+"_frontoparallel.jpg")

```


2.1.2 Result




Figure 5: Warped threebody




Figure 6: Warped palmer

2.1.3 Discussion

The Vision Science (palmer) book's top can bottom edge is curved after the warp action. This is because the surface of this book is slightly concave, therefore the photo taken from side view will cause the edge between the corners are curved as well. The homography (projection transformation) preserve lines inherently, so these curves are kept.

2.1.4 Optional

If the synthetic cover contains only ones, it would act as a white mask. When we apply the inverse warp using the homography that maps the synthetic cover onto the scene image, this white mask would be positioned over the area where the book cover appears in the scene image, making it look like the book cover is entirely white.

2.2 Task6: Stitching Stuff Together

2.2.1 Fill in compute_distance in task6.py

```
1 """
2 Task6 Code
3 """
4 import numpy as np
5 import common
6 from common import save_img, read_img
7 from homography import fit_homography, homography_transform,
8     RANSAC_fit_homography
9 import os
10 import cv2
11
12 def compute_distance(desc1, desc2):
13     """
14     Calculates L2 distance between 2 binary descriptor vectors.
15
16     Input - desc1: Descriptor vector of shape (N,F)
17             desc2: Descriptor vector of shape (M,F)
18
19     Output - dist: a (N,M) L2 distance matrix where dist(i,j)
20                 is the squared Euclidean distance between row i of
21                 desc1 and desc2. You may want to use the distance
22                 calculation trick
23                 
$$||x - y||^2 = ||x||^2 + ||y||^2 - 2x^T y$$

24
25     """
26     X = desc1
27     Y = desc2
28
29     X_norm_sq = np.linalg.norm(X, axis=1, keepdims=True) ** 2
30     Y_norm_sq = np.linalg.norm(Y, axis=1, keepdims=True) ** 2
31     dist = np.sqrt(np.maximum(0, (X_norm_sq + Y_norm_sq.T - 2 * (X @ Y.T))
32 ))
33     return dist
```

2.2.2 Fill in find_matches in task6.py

```
1 def find_matches(desc1, desc2, ratioThreshold):
2     """
3     Calculates the matches between the two sets of keypoint
4     descriptors based on distance and ratio test.
5
6     Input - desc1: Descriptor vector of shape (N,F)
7             desc2: Descriptor vector of shape (M,F)
8             ratioThreshold : maximum acceptable distance ratio between 2
9                             nearest matches
10
11     Output - matches: a list of indices (i,j) 1 <= i <= N, 1 <= j <= M
12                  giving
13                  the matches between desc1 and desc2.
```

```

14         This should be of size (K,2) where K is the number of
15         matches and the row [ii,jj] should appear if desc1[ii,:] and
16         desc2[jj,:] match.
17     """
18     matches = []
19
20     dist = compute_distance(desc1, desc2)
21     idx_smallest_two = np.argsort(dist, axis=1)[: , :2]
22     ratio = np.take_along_axis(dist, idx_smallest_two, axis=1)[: , 0] / np.
take_along_axis(dist, idx_smallest_two, axis=1)[: , 1]
23
24     idx_ii = np.where((ratio < ratioThreshold))[0]
25     idx_jj = idx_smallest_two[idx_ii, 0]
26     matches = np.hstack((idx_ii[:, np.newaxis], idx_jj[:, np.newaxis]))
27     # import pdb; pdb.set_trace()
28     return matches

```

2.2.3 Fill in draw_matches in task6.py

```

1 def draw_matches(img1, img2, kp1, kp2, matches):
2     """
3     Creates an output image where the two source images stacked vertically
4     connecting matching keypoints with a line.
5
6     Input - img1: Input image 1 of shape (H1,W1,3)
7             img2: Input image 2 of shape (H2,W2,3)
8             kp1: Keypoint matrix for image 1 of shape (N,4)
9             kp2: Keypoint matrix for image 2 of shape (M,4)
10            matches: List of matching pairs indices between the 2 sets of
11                    keypoints (K,2)
12
13     Output - Image where 2 input images stacked vertically with lines
14             joining
15             the matched keypoints
16     Hint: see cv2.line
17     """
18     #Hint:
19     #Use common.get_match_points() to extract keypoint locations
20     output = np.vstack((img1, img2))
21     H1, W1, _ = img1.shape
22     kps = common.get_match_points(kp1, kp2, matches)
23     for i in range(kps.shape[0]):
24         p1 = kps[i, :2].astype(int)
25         p2 = (kps[i, 2:] + np.array([0, H1])).astype(int)
26         # print(p1, p2)
27         cv2.line(output, (p1), (p2), (0, 0, 255), 4 )
28     return output

```


2.2.4 Picture of matches



Figure 7: Matches of Lowetag

2.2.5 Fill in RANSAC_fit_homography in homography.py

```
1 def RANSAC_fit_homography(XY, eps=1, nIters=1000):  
2     ...
```

```

3 Perform RANSAC to find the homography transformation
4 matrix which has the most inliers
5
6 Input - XY: an array with size(N,4), each row contains two
7         points in the form [x_i, y_i, x'_i, y'_i] (1,4)
8         eps: threshold distance for inlier calculation
9         nIters: number of iteration for running RANSAC
10 Output - bestH: a (3,3) homography matrix fit to the
11            inliers from the best model.
12
13 Hints:
14 a) Sample without replacement. Otherwise you risk picking a set of
15    points
16    that have a duplicate.
17 b) *Re-fit* the homography after you have found the best inliers
18 '''
19 # bestH, bestCount, bestInliers = np.eye(3), -1, np.zeros((XY.shape
20 [0],))
21 # bestRefit = np.eye(3)
22
23 # Initialize the best homography matrix, inlier count and inlier set
24 bestH = None
25 bestCount = -1
26 bestInliers = None
27
28 for _ in range(nIters):
29     # Step 1: Randomly select 4 pairs of points without replacement
30     indices = np.random.choice(XY.shape[0], 4, replace=False)
31     sample = XY[indices]
32
33     # Step 2: Compute the homography matrix using the provided utility
34     function
35     H = fit_homography(sample)
36
37     # Step 3: Apply homography and determine inliers
38     # Transform source points to destination plane
39     homogenized_src_pts = np.concatenate((XY[:, :2], np.ones((XY.shape
40 [0], 1))), axis=1)
41     transformed_pts = np.dot(H, homogenized_src_pts.T).T
42     transformed_pts /= transformed_pts[:, 2][:, np.newaxis] #
43 Normalize
44
45     # Calculate distances from actual to projected points
46     homogenized_dst_pts = np.concatenate((XY[:, 2:], np.ones((XY.shape
47 [0], 1))), axis=1)
48     distances = np.linalg.norm(homogenized_dst_pts[:, :2] -
49 transformed_pts[:, :2], axis=1)
50
51     # Inliers are points with distance less than epsilon
52     inliers = distances < eps
53     inlier_count = np.sum(inliers)
54
55     # Step 4: Keep track of the best homography with the most inliers
56     if inlier_count > bestCount:

```

```

50         bestCount = inlier_count
51         bestH = H
52         bestInliers = inliers
53
54     # Step 5: Re-fit the homography using all inliers from the best model
    found
55     if bestInliers is not None and bestCount > 4: # More than the minimal
        sample size
56         all_inliers = XY[bestInliers]
57         bestH = fit_homography(all_inliers)
58     else:
59         bestH = np.eye(3) # Fallback to identity matrix if no good model
        is found
60
61     return bestH

```

2.2.6 Fill in make_warped and warp_and_combine in task6.py

```

1 def warp_and_combine(img1, img2, H):
2     """
3     You may want to write a function that merges the two images together
4     given
5     the two images and a homography: once you have the homography you do
6     not
7     need the correspondences; you just need the homography.
8     Writing a function like this is entirely optional, but may reduce the
9     chance
10    of having a bug where your homography estimation and warping code have
11    odd
12    interactions.
13
14    Input - img1: Input image 1 of shape (H1,W1,3)
15           img2: Input image 2 of shape (H2,W2,3)
16           H: homography mapping between them
17    Output - V: stitched image of size (?,?,3); unknown since it depends
18    on H
19    """
20    # Get dimensions of input images
21    h1, w1 = img1.shape[:2]
22    h2, w2 = img2.shape[:2]
23
24    # Corners of img1
25    corners_img1 = np.array([[0, 0], [0, h1], [w1, h1], [w1, 0]], dtype=np
26    .float32).reshape(-1, 1, 2)
27
28    # Corners of img2 transformed by H
29    corners_img2 = np.array([[0, 0], [0, h2], [w2, h2], [w2, 0]], dtype=np
30    .float32).reshape(-1, 1, 2)
31    corners_img2_transformed = cv2.perspectiveTransform(corners_img2, H)
32
33    # Combine the corners
34    all_corners = np.concatenate((corners_img1, corners_img2_transformed),
35    axis=0)

```

```

28
29 # Find the bounding rectangle
30 x_min, y_min = np.intp(np.min(all_corners, axis=0).ravel() - 0.5)
31 x_max, y_max = np.intp(np.max(all_corners, axis=0).ravel() + 0.5)
32
33 # Translation homography
34 translation_dist = [-x_min, -y_min]
35 H_translation = np.array([[1, 0, translation_dist[0]],
36                           [0, 1, translation_dist[1]],
37                           [0, 0, 1]], dtype=np.float32)
38
39 # Warp both images
40 warp_img1 = cv2.warpPerspective(img1, H_translation, (x_max - x_min,
41 y_max - y_min))
42 warp_img2 = cv2.warpPerspective(img2, H_translation.dot(H.astype(np.
43 float32)), (x_max - x_min, y_max - y_min))
44
45 # Create a mask of the combined size for where img1 and warped img2
46 are not zero
47 mask_img1 = np.sum(warp_img1, axis=2) > 0
48 mask_img2 = np.sum(warp_img2, axis=2) > 0
49 mask_overlap = mask_img1 & mask_img2
50 mask_img1_only = mask_img1 & ~mask_overlap
51 mask_img2_only = mask_img2 & ~mask_overlap
52
53 # Initialize the stitched image canvas
54 stitched_img = np.zeros_like(warp_img1)
55
56 # Place each image on the canvas according to the masks
57 stitched_img[mask_img1_only] = warp_img1[mask_img1_only]
58 stitched_img[mask_img2_only] = warp_img2[mask_img2_only]
59
60 # Handle overlapping areas
61 stitched_img[mask_overlap] = warp_img1[mask_overlap] // 2 + warp_img2[
62 mask_overlap] // 2
63
64 return stitched_img

```

```

1 def make_warped(img1, img2):
2     """
3     Take two images and return an image, putting together the full
4     pipeline.
5     You should return an image of the panorama put together.
6
7     Input - img1: Input image 1 of shape (H1,W1,3)
8             img2: Input image 1 of shape (H2,W2,3)
9
10    Output - Final stitched image
11    Be careful about:
12    a) The final image size
13    b) Writing code so that you first estimate H and then merge images
14    with H.
15    The system can fail to work due to either failing to find the
16    homography or

```



```

14     failing to merge things correctly.
15     '''
16
17     kp1, desc1 = common.get_AKAZE(I1)
18     kp2, desc2 = common.get_AKAZE(I2)
19
20     ratio = 0.7
21     matches = find_matches(desc1, desc2, ratio)
22     kps = common.get_match_points(kp1, kp2, matches)
23
24     H = RANSAC_fit_homography(kps, eps= 4, nIters=2000)
25     print(H)
26
27     stitched = warp_and_combine(img2, img1, H)
28
29     return stitched

```

2.2.7 Two panorama figures



Figure 8: Lowetag Panorama



Figure 9: Eynsham Panorama

2.2.8 Include Figures in .zip

Above two figures are submitted in Canvas

3 Augmented Reality on a Budget

3.1 Task7: Augmented Reality on a Budget

3.1.1 Fill in the function `improve_image(scene,template,transfer)` in `task7.py`

```

1  """
2  Task 7 Code
3  """
4  import numpy as np
5  import common
6  from common import save_img, read_img
7  from homography import homography_transform, RANSAC_fit_homography
8  import cv2
9  import os
10
11  from task6 import *
12
13  def task7_warp_and_combine(img1, img2, H):
14      """
15      You may want to write a function that merges the two images together
16      given
17      the two images and a homography: once you have the homography you do
18      not
19      need the correspondences; you just need the homography.

```

```

18     Writing a function like this is entirely optional, but may reduce the
19     chance
20     of having a bug where your homography estimation and warping code have
21     odd
22     interactions.
23
24     Input - img1: Input image 1 of shape (H1,W1,3)
25             img2: Input image 2 of shape (H2,W2,3)
26             H: homography mapping between them
27     Output - V: stitched image of size (?,?,3); unknown since it depends
28     on H
29             but make sure in V, for pixels covered by both img1 and
30     warped img2,
31             you see only img2
32     '''
33     # Warp img2 onto img1's plane
34     warp_img2 = cv2.warpPerspective(img2, H, (img1.shape[1], img1.shape
35     [0]))
36     # Create mask of where the warped image is non-zero
37     mask = (warp_img2.sum(-1) > 0)
38     # Initialize output image
39     V = img1.copy()
40     # Place img2 on the masked regions of img1
41     V[mask] = warp_img2[mask]
42
43     return V
44
45 def improve_image(scene, template, transfer):
46     '''
47     Detect template image in the scene image and replace it with transfer
48     image.
49
50     Input - scene: image (H,W,3)
51             template: image (K,K,3)
52             transfer: image (L,L,3)
53     Output - augment: the image with
54
55     Hints:
56     a) You may assume that the template and transfer are both squares.
57     b) This will work better if you find a nearest neighbor for every
58     template
59         keypoint as opposed to the opposite, but be careful about
60     directions of the
61         estimated homography and warping!
62     '''
63     # augment = None
64     # Resize transfer image to the template's size
65     transfer = cv2.resize(transfer, (template.shape[1], template.shape[0])
66     )
67
68     kp1, desc1 = common.get_AKAZE(template)
69     kp2, desc2 = common.get_AKAZE(scene)
70
71     ratio = 0.7

```

```

63 matches = find_matches(desc1, desc2, ratio)
64 kps = common.get_match_points(kp1, kp2, matches)
65
66 H = RANSAC_fit_homography(kps, eps= 4, nIters=2000)
67
68 augment = task7_warp_and_combine(scene, transfer, H)
69
70 return augment
71
72 if __name__ == "__main__":
73     # Task 7
74     scene_img_path = 'task7/scenes/lacroix/scene.jpg'
75     template_img_path = 'task7/scenes/lacroix/template.png'
76     transfer_img_path = 'task7/seals/monk.png'
77     # scene_img_path = 'task7/scenes/bbb/scene.jpg'
78     # template_img_path = 'task7/scenes/bbb/template.png'
79     # transfer_img_path = 'task7/seals/um.png'
80
81     scene = read_img(scene_img_path)
82     template = read_img(template_img_path)
83     transfer = read_img(transfer_img_path)
84
85     improved_image = improve_image(scene, template, transfer)
86
87     save_img(improved_image, f'improved_lacroix.jpg' )

```

3.1.2 Result



Figure 10: BBB Scene



Figure 11: BBB Template



Figure 12: UM logo To Transfer



Figure 13: Augmented BBB

Submitted by Wensong Hu on March 3, 2024.