

EECS 442 Homework 4: Fashion-MNIST Classification

In this part, you will implement and train Convolutional Neural Networks (ConvNets) in PyTorch to classify images. Unlike HW4 Section 1, backpropagation is automatically inferred by PyTorch, so you only need to write code for the forward pass.

Before we start, please put your name and UMID in following format

: Firstname LASTNAME, #00000000 // e.g.) David FOUHEY, #12345678

Your Answer:

Wensong HU #24908654

Setup

```
In [2]: # Run the command in the terminal if it failed on local Jupyter Notebook, remove  
# !pip install torchsummary
```

```
In [3]: import numpy as np  
import matplotlib.pyplot as plt  
from tqdm import tqdm # Displays a progress bar  
  
import torch  
from torch import nn  
from torch import optim  
import torch.nn.functional as F  
from torchsummary import summary  
from torchvision import datasets, transforms  
from torch.utils.data import Dataset, Subset, DataLoader, random_split
```

```
In [4]: if torch.cuda.is_available():  
    print("Using the GPU. You are good to go!")  
    device = 'cuda'  
else:  
    print("Using the CPU. Overall speed may be slowed down")  
    device = 'cpu'
```

Using the GPU. You are good to go!

Loading Dataset

The dataset we use is Fashion-MNIST dataset, which is available at <https://github.com/zalandoresearch/fashion-mnist> and in `torchvision.datasets`. Fashion-MNIST has 10 classes, 60000 training+validation images (we have splitted it to have 50000 training images and 10000 validation images, but you can change the numbers), and 10000 test images.

```
In [5]: # Load the dataset and train, val, test splits
print("Loading datasets...")
# Transform from [0,255] uint8 to [0,1] float,
# then normalize to zero mean and unit variance
FASHION_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize([0.2859], [0.3530])
])
FASHION_trainval = datasets.FashionMNIST('.', download=True, train=True,
                                         transform=FASHION_transform)
FASHION_train = Subset(FASHION_trainval, range(50000))
FASHION_val = Subset(FASHION_trainval, range(50000, 60000))
FASHION_test = datasets.FashionMNIST('.', download=True, train=False,
                                      transform=FASHION_transform)
print("Done!")
```

Loading datasets...
Done!

Now, we will create the dataloader for train, val and test dataset. You are free to experiment with different batch sizes.

```
In [6]: # Create dataloaders
#####
# TODO: Experiment with different batch sizes                                     #
#####
batch_size = 64
#####
#                                     END OF YOUR CODE                           #
#####
trainloader = DataLoader(FASHION_train, batch_size=batch_size, shuffle=True)
valloader = DataLoader(FASHION_val, batch_size=batch_size, shuffle=True)
testloader = DataLoader(FASHION_test, batch_size=batch_size, shuffle=True)
```

Model

Initialize your model and experiment with with different optimizers, parameters (such as learning rate) and number of epochs.

```
In [7]: class Network(nn.Module):
    def __init__(self):
        super().__init__()
        #####
        # TODO: Design your own network, define layers here.
        # Here We provide a sample of two-layer fc network from HW4 Part3.
        # Your solution, however, should contain convolutional layers.
        # Refer to PyTorch documentations of torch.nn to pick your layers.
        # (https://pytorch.org/docs/stable/nn.html)
        # Some common choices: Linear, Conv2d, ReLU, MaxPool2d, AvgPool2d, Drop
        # If you have many layers, use nn.Sequential() to simplify your code
        #####
        # stem: 3*28*28 -> 64 * 14* 14
        self.stem = torch.nn.Sequential(torch.nn.Conv2d(1, 64, kernel_size=7,
                                                         torch.nn.BatchNorm2d(64),
                                                         torch.nn.MaxPool2d(kernel_size=3, stric
        # stage1: 64 * 14 * 14 -> 64 * 14 * 14
```

```

self.resblock1 = torch.nn.Sequential(torch.nn.Conv2d(64, 64, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(64),
                                      torch.nn.ReLU(inplace=True),
                                      torch.nn.Conv2d(64, 64, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(64))
self.resblock2 = torch.nn.Sequential(torch.nn.Conv2d(64, 64, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(64),
                                      torch.nn.ReLU(inplace=True),
                                      torch.nn.Conv2d(64, 64, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(64))

# stage2: 64 * 14 * 14 -> 128 * 7 * 7
self.resblock3 = torch.nn.Sequential(torch.nn.Conv2d(64, 128, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(128),
                                      torch.nn.ReLU(inplace=True),
                                      torch.nn.Conv2d(128, 128, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(128))
self.resblock4 = torch.nn.Sequential(torch.nn.Conv2d(128, 128, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(128),
                                      torch.nn.ReLU(inplace=True),
                                      torch.nn.Conv2d(128, 128, kernel_size=3, padding=1),
                                      torch.nn.BatchNorm2d(128))

# fc layer:
self.pool = torch.nn.AvgPool2d(kernel_size=7)
self.fc = torch.nn.Sequential(torch.nn.Linear(128, 512),
                              torch.nn.ReLU(inplace=True),
                              torch.nn.Linear(512, 10))

# downsample
self.projection = torch.nn.Sequential(torch.nn.Conv2d(64, 128, kernel_size=3, padding=1),
                                       torch.nn.BatchNorm2d(128))

#ReLU
self.relu = torch.nn.ReLU(inplace=True)
#####
#                               END OF YOUR CODE
#####

def forward(self, x):
#####
# TODO: Design your own network, implement forward pass here
#####
# print(x.shape)
N, C, H, W = x.shape

# stem
x = self.stem(x)
# print(x.shape)

# stage1
x_res = torch.clone(x)
x = self.resblock1(x)
x += x_res
x = self.relu(x)
# print(x.shape)

x_res = torch.clone(x)
x = self.resblock2(x)
x += x_res
x = self.relu(x)
# print(x.shape)

```

```

# stage2
x_res = self.projection(x)
x = self.resblock3(x)
x += x_res
x = self.relu(x)
# print(x.shape)

x_res = torch.clone(x)
x = self.resblock4(x)
x += x_res
x = self.relu(x)
# print(x.shape)

# fc layer
x = self.pool(x)
# print(x.shape)
x = torch.flatten(x, start_dim=1)
# print(x.shape)
x = self.fc(x)

return x
#####
#                                     END OF YOUR CODE
#####

model = Network().to(device)
criterion = nn.CrossEntropyLoss() # Specify the loss layer
print('Your network:')
print(summary(model, (1,28,28), device=device)) # visualize your model

#####
# TODO: Modify the lines below to experiment with different optimizers,      #
# parameters (such as learning rate) and number of epochs.                  #
#####
# Set up optimization hyperparameters
learning_rate, weight_decay, num_epoch = 1e-4, 0.0, 5
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate, weight_decay=weight_decay)
#####
#                                     END OF YOUR CODE
#####

```

Your network:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 28, 28]	3,200
BatchNorm2d-2	[-1, 64, 28, 28]	128
MaxPool2d-3	[-1, 64, 14, 14]	0
Conv2d-4	[-1, 64, 14, 14]	36,928
BatchNorm2d-5	[-1, 64, 14, 14]	128
ReLU-6	[-1, 64, 14, 14]	0
Conv2d-7	[-1, 64, 14, 14]	36,928
BatchNorm2d-8	[-1, 64, 14, 14]	128
ReLU-9	[-1, 64, 14, 14]	0
Conv2d-10	[-1, 64, 14, 14]	36,928
BatchNorm2d-11	[-1, 64, 14, 14]	128
ReLU-12	[-1, 64, 14, 14]	0
Conv2d-13	[-1, 64, 14, 14]	36,928
BatchNorm2d-14	[-1, 64, 14, 14]	128
ReLU-15	[-1, 64, 14, 14]	0
Conv2d-16	[-1, 128, 7, 7]	8,320
BatchNorm2d-17	[-1, 128, 7, 7]	256
Conv2d-18	[-1, 128, 7, 7]	73,856
BatchNorm2d-19	[-1, 128, 7, 7]	256
ReLU-20	[-1, 128, 7, 7]	0
Conv2d-21	[-1, 128, 7, 7]	147,584
BatchNorm2d-22	[-1, 128, 7, 7]	256
ReLU-23	[-1, 128, 7, 7]	0
Conv2d-24	[-1, 128, 7, 7]	147,584
BatchNorm2d-25	[-1, 128, 7, 7]	256
ReLU-26	[-1, 128, 7, 7]	0
Conv2d-27	[-1, 128, 7, 7]	147,584
BatchNorm2d-28	[-1, 128, 7, 7]	256
ReLU-29	[-1, 128, 7, 7]	0
AvgPool2d-30	[-1, 128, 1, 1]	0
Linear-31	[-1, 512]	66,048
ReLU-32	[-1, 512]	0
Linear-33	[-1, 10]	5,130

Total params: 748,938
 Trainable params: 748,938
 Non-trainable params: 0

Input size (MB): 0.00
 Forward/backward pass size (MB): 2.69
 Params size (MB): 2.86
 Estimated Total Size (MB): 5.55

None

Run the cell below to start your training, we expect you to achieve over **85%** on the test set. A valid solution that meet the requirement take no more than **10 minutes** on normal PC Intel core CPU setting. If your solution takes too long to train, try to simplify your model or reduce the number of epochs.

```
In [8]: %time
def train(model, trainloader, valloader, num_epoch=10): # Train the model
    print("Start training...")
    trn_loss_hist = []
```

```

trn_acc_hist = []
val_acc_hist = []
model.train() # Set the model to training mode
for i in range(num_epoch):
    running_loss = []
    print('-----Epoch = %d-----' % (i+1))
    for batch, label in tqdm(trainloader):
        batch = batch.to(device)
        label = label.to(device)
        optimizer.zero_grad() # Clear gradients from the previous iteration
        # This will call Network.forward() that you implement
        pred = model(batch)
        loss = criterion(pred, label) # Calculate the loss
        running_loss.append(loss.item())
        loss.backward() # Backprop gradients to all tensors in the network
        optimizer.step() # Update trainable weights
    print("\n Epoch {} loss:{}".format(i+1, np.mean(running_loss)))

    # Keep track of training loss, accuracy, and validation loss
    trn_loss_hist.append(np.mean(running_loss))
    trn_acc_hist.append(evaluate(model, trainloader))
    print("\n Evaluate on validation set...")
    val_acc_hist.append(evaluate(model, valloader))
print("Done!")
return trn_loss_hist, trn_acc_hist, val_acc_hist

def evaluate(model, loader): # Evaluate accuracy on validation / test set
    model.eval() # Set the model to evaluation mode
    correct = 0
    with torch.no_grad(): # Do not calculate gradient to speed up computation
        for batch, label in tqdm(loader):
            batch = batch.to(device)
            label = label.to(device)
            pred = model(batch)
            correct += (torch.argmax(pred, dim=1) == label).sum().item()
    acc = correct/len(loader.dataset)
    print("\n Evaluation accuracy: {}".format(acc))
    return acc

trn_loss_hist, trn_acc_hist, val_acc_hist = train(model, trainloader,
                                                  valloader, num_epoch)

#####
# TODO: Note down the evaluation accuracy on test set #
#####
print("\n Evaluate on test set")
evaluate(model, testloader)
#5: 64, 1e-4, 0.95, 5: 0.7843
#6: 64, 1e-4, 0.98, 5: 0.7851
#7: 64, 1e-4, 0.00, 5: 0.8917

```

Start training...

-----Epoch = 1-----

0%| | 0/782 [00:00<?, ?it/s]100%|██████████| 782/782 [00:13<00:00, 59.46it/s]

Epoch 1 loss:0.5725756988424779

100%|██████████| 782/782 [00:07<00:00, 99.74it/s]

Evaluation accuracy: 0.83086

Evaluate on validation set...

100%|██████████| 157/157 [00:01<00:00, 96.58it/s]

Evaluation accuracy: 0.8212

-----Epoch = 2-----

100%|██████████| 782/782 [00:12<00:00, 61.51it/s]

Epoch 2 loss:0.38106703539105025

100%|██████████| 782/782 [00:07<00:00, 101.50it/s]

Evaluation accuracy: 0.89394

Evaluate on validation set...

100%|██████████| 157/157 [00:01<00:00, 98.84it/s]

Evaluation accuracy: 0.8852

-----Epoch = 3-----

100%|██████████| 782/782 [00:12<00:00, 61.09it/s]

Epoch 3 loss:0.31754977807707496

100%|██████████| 782/782 [00:07<00:00, 102.77it/s]

Evaluation accuracy: 0.89466

Evaluate on validation set...

100%|██████████| 157/157 [00:01<00:00, 97.87it/s]

Evaluation accuracy: 0.8833

-----Epoch = 4-----

100%|██████████| 782/782 [00:12<00:00, 61.65it/s]

Epoch 4 loss:0.2882935323983507

100%|██████████| 782/782 [00:07<00:00, 101.61it/s]

Evaluation accuracy: 0.89676

Evaluate on validation set...

100%|██████████| 157/157 [00:01<00:00, 101.90it/s]

Evaluation accuracy: 0.8843

-----Epoch = 5-----

100%|██████████| 782/782 [00:12<00:00, 61.54it/s]

Epoch 5 loss:0.2685324148086788

100%|██████████| 782/782 [00:07<00:00, 99.73it/s]

Evaluation accuracy: 0.9109

Evaluate on validation set...

100%|██████████| 157/157 [00:01<00:00, 101.33it/s]

Evaluation accuracy: 0.8966

Done!

Evaluate on test set

100%|██████████| 157/157 [00:01<00:00, 103.19it/s]

Evaluation accuracy: 0.8917

CPU times: user 1min 54s, sys: 512 ms, total: 1min 54s

Wall time: 1min 52s

Out[8]: 0.8917

Once your training is complete, run the cell below to visualize the training and validation accuracies across iterations.

```
In [9]: #####  
# TODO: Submit the accuracy plot #  
#####  
# visualize the training / validation accuracies  
x = np.arange(num_epoch)  
# train/val accuracies for MiniVGG  
plt.figure()  
plt.plot(x, trn_acc_hist)  
plt.plot(x, val_acc_hist)  
plt.legend(['Training', 'Validation'])  
plt.xticks(x)  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.title('fashion MNIST Classification')  
plt.gcf().set_size_inches(10, 5)  
plt.savefig('part1.png', dpi=300)  
plt.show()
```

