

HW5 — Generative Models

1 Section 1: Pix2Pix

1.1 Task 1: Dataloading

Code submitted to Canvas. Full notebook see appendix.

```
1 class Edges2Image(Dataset):
2     def __init__(self, root_dir, split='train', transform=None):
3         """
4         Args:
5             root_dir: the directory of the dataset
6             split: "train" or "val"
7             transform: pytorch transformations.
8         """
9
10        self.transform = transform
11
12        self.files = glob.glob(os.path.join(root_dir, split, '*.jpg'))
13
14    def __len__(self):
15        return len(self.files)
16
17    def __getitem__(self, idx):
18        img = Image.open(self.files[idx])
19        img = np.asarray(img)
20        if self.transform:
21            img = self.transform(img)
22        return img
23
24    transform = transforms.Compose([
25        transforms.ToTensor(),
26        transforms.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
27    ])
28
29    #
30    # #####
31    # TODO: Construct the dataloader
32    #
33    # For the train_loader, please use a batch size of 4 and set shuffle True
34    #
```

```

32 # For the val_loader, please use a batch size of 5 and set shuffle False
    #
33 # Hint: You'll need to create instances of the class above, name them as
    #
34 # tr_dt and te_dt. The dataloaders should be named as train_loader and
    #
35 # test_loader. You also need to include transform in your class
    #
36 #instances
    #
37 #
    #####
38
39 tr_dt = Edges2Image(root_dir="mini-edges2shoes", split='train', transform=
    transform)
40 te_dt = Edges2Image(root_dir="mini-edges2shoes", split='val', transform=
    transform)
41
42 train_loader = DataLoader(tr_dt, batch_size=4, shuffle=True)
43 test_loader = DataLoader(te_dt, batch_size=5, shuffle=True)
44
45 #
    #####
46 #
    #####
47 #
    #####
48
49 # Make sure that you have 1,000 training images and 100 testing images
    before moving on
50 print('Number of training images {}, number of testing images {}'.format(
    len(tr_dt), len(te_dt)))

```

1.2 Task 2: Training Pix2Pix

1.2.1 Codes

Code submitted to Canvas. Full notebook see appendix.

```

1 # Hint: you could use following loss to complete following function
2 BCE_loss = nn.BCELoss().cuda()
3 L1_loss = nn.L1Loss().cuda()
4
5 def train(G, D, num_epochs = 20):
6     hist_D_losses = []
7     hist_G_losses = []
8     hist_G_L1_losses = []
9     #
    #####

```

```

10 # TODO: Add Adam optimizer to generator and discriminator
11 #
12 # You will use lr=0.0002, beta=0.5, beta2=0.999
13 #
14 #####
15 G_optimizer = optim.Adam(G.parameters(), lr=2e-4, betas=(0.5, 0.999))
16 D_optimizer = optim.Adam(D.parameters(), lr=2e-4, betas=(0.5, 0.999))
17 #
18 #####
19 #
20 #
21 #
22 #####
23
24 print('training start!')
25 start_time = time.time()
26 for epoch in range(num_epochs):
27     print('Start training epoch %d' % (epoch + 1))
28     D_losses = []
29     G_losses = []
30     epoch_start_time = time.time()
31     num_iter = 0
32     for x_ in train_loader:
33         y_ = x_[:, :, :, img_size:]
34         x_ = x_[:, :, :, 0:img_size]
35
36         x_, y_ = x_.cuda(), y_.cuda()
37         #
38         #####
39
40         # TODO: Implement training code for the discriminator.
41         #
42         # Recall that the loss is the mean of the loss for real images and
43         # fake images, and made by some calculations with zeros and ones
44         #
45         # We have defined the BCE_loss, which you might would like to use.
46         #
47         #
48         #
49         # NOTE: While training the Discriminator, the output of the
50         # generator must be detached from the computational graph. Refer to the method
51         # torch.Tensor.detach()
52         #
53         #
54         #####

```

```

44     N = x_.shape[0]
45     # Generate data
46     fake_data = G.forward(x_).detach()
47
48
49     #1. Train the discriminator
50     # D real data BCE loss
51     D_real_preds = D.forward(torch.cat((x_, y_), dim=1))
52     D_y_real = torch.ones_like(D_real_preds)
53     # D_real_loss = torch.sum(torch.log(D_real_preds))
54     D_real_loss = BCE_loss(D_real_preds, D_y_real)
55
56     # D fake data BCE loss
57     D_fake_preds = D.forward(torch.cat((x_, fake_data), dim=1))
58     D_y_fake = torch.zeros_like(D_fake_preds)
59     # D_fake_loss = torch.sum(torch.log(1 - D_fake_preds))
60     D_fake_loss = BCE_loss(D_fake_preds, D_y_fake)
61
62     # D loss
63     loss_D = D_real_loss + D_fake_loss
64
65     # Train D
66     D_optimizer.zero_grad()
67     loss_D.backward()
68     D_optimizer.step()
69
70     #
71     #####
72
73     #                                     END OF YOUR CODE
74     #
75     #
76     #####
77
78     # TODO: Implement training code for the Generator.
79     #
80     #
81     #####
82
83     # 1. Train the generator
84     # 2. Append the losses to the lists 'hist_G_L1_losses' and '
85     hist_G_losses'
86     # (Only append the data to the list, not the complete tensor, refer
87     # torch.Tensor.item()).
88
89     # Generate data
90     fake_data = G.forward(x_)

```

```

86     # 1. Train the generator
87     # G BCE loss
88     G_fake_preds = D.forward(torch.cat((x_, fake_data), dim=1))
89     G_y_fake = torch.zeros_like(G_fake_preds)
90     G_bce_loss = BCE_loss(G_fake_preds, G_y_fake)
91
92     # G l1 loss
93     G_l1_loss = L1_loss(fake_data, y_)
94
95     # G loss
96     lamb = 100
97     loss_G = G_bce_loss + lamb * G_l1_loss
98
99     # Train G
100    G_optimizer.zero_grad()
101    loss_G.backward()
102    G_optimizer.step()
103
104    # 2. Append the losses to the lists 'hist_G_L1_losses' and '
hist_D_losses'
105    # (Only append the data to the list, not the complete tensor, refer
106    # torch.Tensor.item()).
107    hist_G_losses.append(G_bce_loss.detach().item())
108    hist_G_L1_losses.append(G_l1_loss.detach().item())
109    #
#####
110
111    #                                     END OF YOUR CODE
112    #
113    #
#####
114
115    D_losses.append(loss_D.detach().item())
116    hist_D_losses.append(loss_D.detach().item())
117    G_losses.append(loss_G)
118    num_iter += 1
119
120    epoch_end_time = time.time()
121    per_epoch_ptime = epoch_end_time - epoch_start_time
122
123    print('[%d/%d] - using time: %.2f seconds' % ((epoch + 1), num_epochs,
per_epoch_ptime))
124    print('loss of discriminator D: %.3f' % (torch.mean(torch.FloatTensor(
D_losses))))
125    print('loss of generator G: %.3f' % (torch.mean(torch.FloatTensor(
G_losses))))
126    if epoch == 0 or (epoch + 1) % 5 == 0:
127        with torch.no_grad():
128            show_result(G, fixed_x_, fixed_y_, (epoch+1))
129
130    end_time = time.time()

```

```

131 total_ptime = end_time - start_time
132
133 return hist_D_losses, hist_G_losses, hist_G_L1_losses

```

1.2.2 Report Results



Figure 1: Result Visualization after 20 training epoches

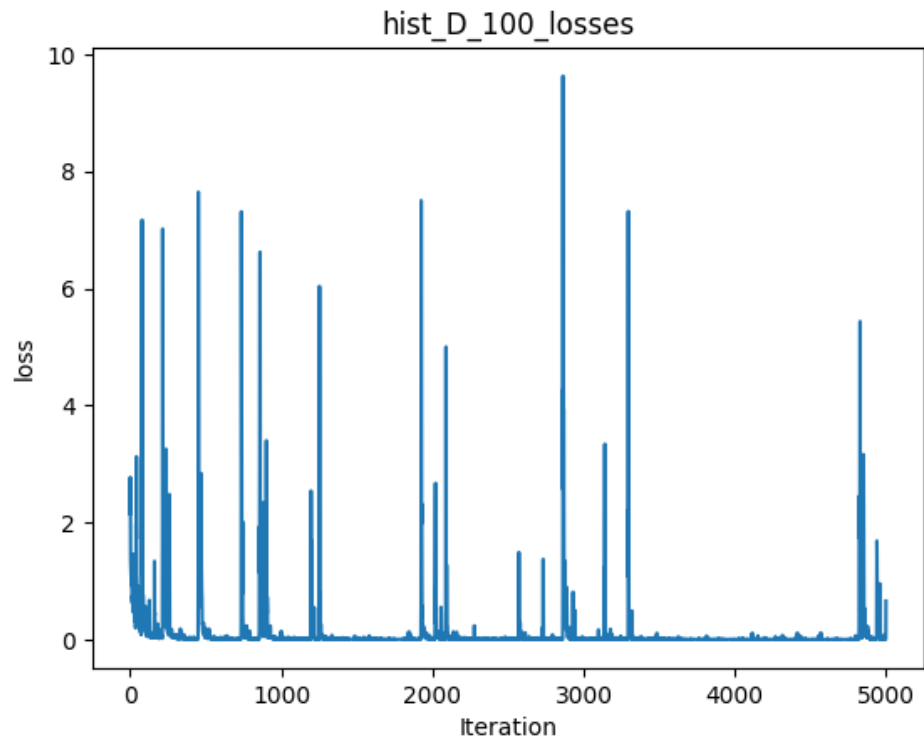


Figure 2: Discriminator BCE loss

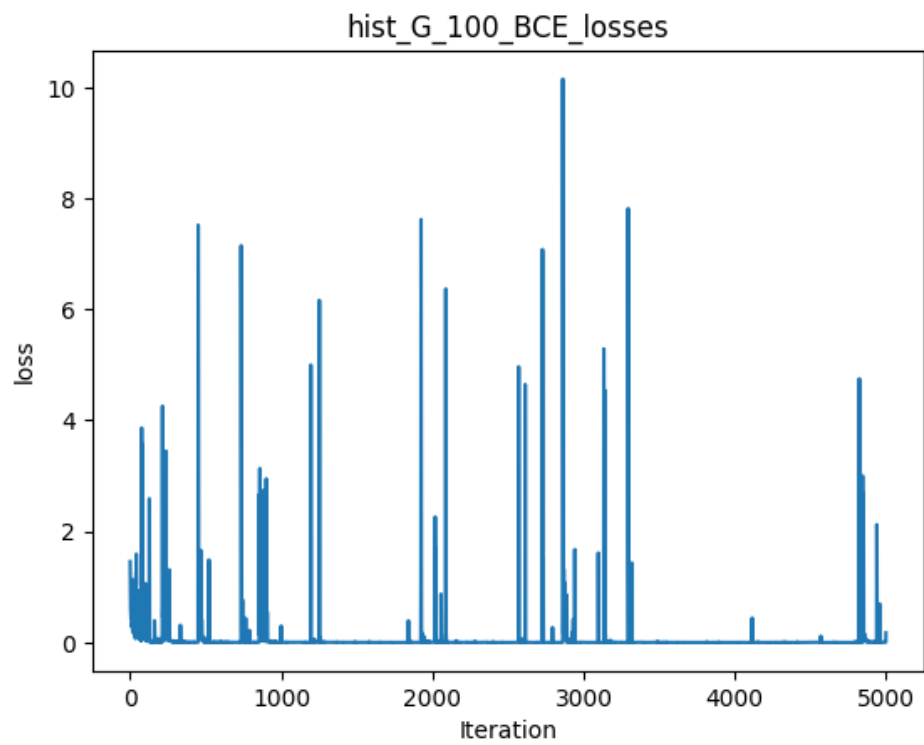


Figure 3: Generator BCE loss

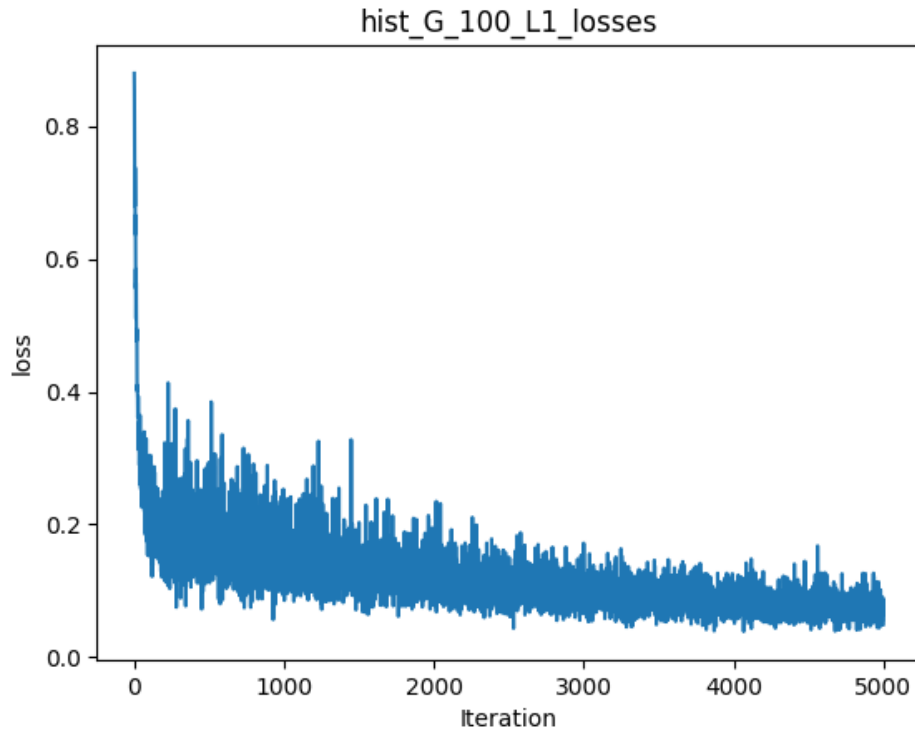


Figure 4: Generator L1 loss

2 Section 2: Diffusion Models

2.1 Task 3: Unconditional Sampling using DDPM: get_name_beta_schedule

Code submitted to Canvas.

```

1 def get_named_beta_schedule(schedule_name, num_diffusion_timesteps,
2   beta_min=0.0001, beta_max=0.02):
3   """
4   Get a pre-defined beta schedule for the given name.
5
6   Args:
7     schedule_name: str, name of the variance schedule, 'linear' or '
8     cosine'
9     num_diffusion_timesteps: int, number of the entire diffusion
10    timesteps
11    beta_min: float, minimum value of beta
12    beta_max: float, maximum value of beta
13
14  Returns:
15    betas: np.ndarray, a 1-d array of size num_diffusion_timesteps,
16    contains all the beta for each timestep
17
18    """
19    betas = None
20    if schedule_name == "linear":
21        ##### START TODO #####

```



```

18     # Implement the linear schedule
19     # Uniformly divide the [beta_min, beta_max) to
num_diffusion_timesteps values.
20     betas = np.linspace(beta_min, beta_max, num_diffusion_timesteps)
21     ##### END TODO #####
22
23
24     elif schedule_name == "cosine":
25         ##### START TODO #####
26         # Implement the cosine schedule
27         # Assume s = 0.008 and beta_clip=0.999
28         s = 0.008
29         beta_clip = 0.999
30
31         betas = np.zeros(num_diffusion_timesteps)
32         f_0 = 0.0
33         alpha_bar_tminus1 = 1.0
34         for t in range(num_diffusion_timesteps):
35             if t == 0:
36                 f_0 = (np.cos((0/num_diffusion_timesteps + s) / (1 + s) *
np.pi / 2)) ** 2
37                 # alpha_bar_tminus1 = 1
38
39             else:
40                 f_t = (np.cos((t/num_diffusion_timesteps + s) / (1 + s) *
np.pi / 2)) ** 2
41                 alpha_bar_t = f_t / f_0
42                 # import pdb; pdb.set_trace()
43                 betas[t-1] = np.clip((1 - alpha_bar_t / alpha_bar_tminus1)
, a_min=0, a_max=beta_clip)
44                 alpha_bar_tminus1 = alpha_bar_t
45                 # import pdb; pdb.set_trace()
46                 betas[-1] = beta_clip
47
48             ##### End TODO #####
49         else:
50             raise NotImplementedError(f"unknown beta schedule: {schedule_name}
")
51
52     return betas

```

2.2 Task 4: Unconditional Sampling using DDPM: DDPM

2.2.1 Code

Code submitted to Canvas:

```

1 @register_sampler("ddpm")
2 class DDPMDiffusion:
3     def __init__(self,
4                 betas,
5                 dynamic_threshold,
6                 clip_denoised,
7                 rescale_timesteps,

```

```

8         **kwargs
9     ):
10
11     # use float64 for accuracy.
12     betas = np.array(betas, dtype=np.float64)
13     self.betas = betas
14     assert self.betas.ndim == 1, "betas must be 1-D"
15     assert (0 < self.betas).all() and (self.betas <=1).all(), "betas
must be in (0..1]"
16
17     self.num_timesteps = int(self.betas.shape[0])
18     self.rescale_timesteps = rescale_timesteps
19
20     ##### START TODO #####
21     # Calculate the values of alpha
22     # Also we will need the cumulated product of alpha.
23     # And during sampling we need the value of cumulated product of
alpha from
24     # previous or next timestep.
25     self.alphas = 1 - betas
26     self.alphas_cumprod = np.cumprod(self.alphas) # cumulated
product of alphas
27     self.alphas_cumprod_prev = np.concatenate((np.array([1]), self.
alphas_cumprod[:-1])) # first T-1 elements of alphas_cumprod, append
1.0 at the begining, to make its length T
28     self.alphas_cumprod_next = np.concatenate((self.alphas_cumprod
[1:], np.array([0]))) # last T-1 elements of alphas_cumprod, append 0.0
at the end, to make its length T
29
30     ##### END TODO #####
31     assert self.alphas_cumprod_prev.shape == (self.num_timesteps,)
32
33     self.mean_processor = EpsilonXMeanProcessor(betas=betas,
dynamic_threshold=
34     dynamic_threshold,
35     clip_denoised=
clip_denoised)
36
37     self.var_processor = LearnedRangeVarianceProcessor(betas=betas)
38
39     def p_sample_loop(self,
40         model,
41         x_start,
42         record,
43         save_root,
44         measurement=None,
45         measurement_cond_fn=None,
46         uncond=False):
47
48         """
49         The function used for sampling from noise.
50
51         Args:
52             model: nn.Module, the pretrained model that is used to predict
the score and variance

```

```

52         x_start: torch.Tensor, random noise input
53         measurement: torch.Tensor, our corrupted observation
54         measurement_cond_fn: conditional function used to perform
conditional sampling, is None for unconditional sampling
55         record: Bool, save intermediate results if True
56         save_root: str, root of the directory to save the results
57         uncond: Bool, perform unconditional sampling if True, else
perform conditional sampling
58         """
59         if not uncond:
60             assert measurement is not None and measurement_cond_fn is not
None, \
61                 "measurement and measurement conditional function is
required for conditional sampling"
62
63         img = x_start # start from random noise
64         device = x_start.device
65
66         ##### Start TODO #####
67         # Implement the sample loop
68         # Call p_sample for every iteration
69         # It requires only one line of code implementation here
70
71         pbar = tqdm(list(range(self.num_timesteps))[:-1])
72         for idx in pbar:
73             time = torch.tensor([idx] * img.shape[0], device=device)
74
75             img = self.p_sample(model=model, x=img, t=time)['x_t_minus_1']
76
77             img = img.detach_()
78
79             if record:
80                 if idx % 10 == 0:
81                     file_path = os.path.join(save_root, f"progress/x_{str(
idx).zfill(4)}.png")
82                     plt.imsave(file_path, clear_color(img))
83
84         return img
85
86     def p_sample(self, model, x, t):
87         """
88         Posterior sampling process, when given the model, x_t and timestep
t, it returns predicted
89         x_0 and x_t_minus_1
90
91         We have already provide you with the function to get the log of
the variance.
92         Use self.var_processor.get_variance(var_values, t), where
var_values is
93         the 3:6 channels of the direct output of the model.
94         example usage: log_variance = self.var_processor.get_variance(
var_values, t)
95

```

```

96     You can also use the helper function extract_and_expand() to
    extract the value
97     corresponding to timestep and expand it to the same size as the
    target for broadcast.
98     example usage: coef1 = extract_and_expand(self.
    posterior_mean_coef1, t, x_start)
99
100     Args:
101         model: nn.Module, the UNet model, you can call model(x, t) to
    get the output tensor with size (B, 6, H, W)
102         x: torch.Tensor, shape (1, 3, H, W), x_t
103         t: torch.Tenosr, shape (1,), timestep
104
105     Returns:
106         output_dict: dict, contains predicted x_t_minus_1 and x_0
107         """
108         #####Start TODO#####
109         ##### Get the predicted score and variance of the pretrained model
    #####
110         model_output = model.forward(x, t)
111         pred_noise = model_output[:, :3, :, :]
112         var_values = model_output[:, 3:, :, :]
113         ##### End TODO #####
114
115         log_variance = self.var_processor.get_variance(var_values, t)    #
    get the log of variance
116
117         ##### Start TODO #####
118         ##### get predicted x_0 and x_t_minus_1 #####
119         ##### don't forget to add noise for all the steps, except for the
    last one #####
120         if t > 1:
121             z = torch.randn(x.shape, dtype=x.dtype, device=x.device)
122         else:
123             z = torch.zeros_like(x, device=x.device)
124         # import pdb; pdb.set_trace()
125         alpha = extract_and_expand(self.alphas, t, x)
126         alpha_bar = extract_and_expand(self.alphas_cumprod, t, x)
127         x_t_minus_1 = (1 / torch.sqrt(alpha)) * (x - ((1 - alpha) / torch.
    sqrt(1 - alpha_bar)) * pred_noise) + torch.sqrt(torch.exp(log_variance)
    ) * z
128
129         ##### End TODO #####
130
131         assert x_t_minus_1.shape == log_variance.shape == x.shape
132
133         output_dict = {'x_t_minus_1': x_t_minus_1}
134         return output_dict

```

2.2.2 Result



Figure 5: DDPM Sampling Result

2.3 Task 5: Unconditional Sampling using DDIM

2.4 Code

Code submitted to Canvas.

```
1 @register_sampler("ddim")
2 class DDIMDiffusion(DDPMDiffusion):
3
4     def __init__(self, use_timesteps, **kwargs):
5
6         self.timestep_map = []
7         self.original_num_steps = len(kwargs["betas"])
8
9         base_alphas_cumprod = DDPMDiffusion(**kwargs).alphas_cumprod #
10 pylint: disable=missing-kwargs
11         last_alpha_cumprod = 1.0
12         new_betas = []
13         self.use_timesteps = set(use_timesteps)
14
15         for i, alpha_cumprod in enumerate(base_alphas_cumprod):
16             if i in self.use_timesteps:
17                 new_betas.append(1 - alpha_cumprod / last_alpha_cumprod)
18                 last_alpha_cumprod = alpha_cumprod
19                 self.timestep_map.append(i)
20         kwargs["betas"] = np.array(new_betas)
21         super().__init__(**kwargs)
22
```

```

23     def _scale_timesteps(self, t):
24         if self.rescale_timesteps:
25             return t.float() * (self.original_num_steps / self.
num_timesteps)
26         return t
27
28     def p_sample(self, model, x, t, eta=0.0):
29         #####
30         ##### TODO #####
31         ##### Get the predicted score and variance of the pretrained model
#####
32         ##### Don't forget to use _scale_timesteps to scale the timestep
for calling the model prediction.
33         ##### You don't need to scale the timestep for further
computations of x_t_minus_1.
34         ##### NOTE: Since this version of the model learns the variance
along with the score function,
35         ##### the output of the model would have double the number of
channels as that of the input.
36         ##### So assign the predicted score and variance values to the
variables below. Refer to
37         ##### torch.split method.
38         #####
39         ##### Start TODO #####
40         model_output = model.forward(x, self._scale_timesteps(t))
41         # import pdb; pdb.set_trace()
42         pred_noise, var_values = torch.split(model_output, 3, dim=1)
43         ##### End TODO #####
44
45         model_mean, pred_xstart = self.mean_processor.get_mean_and_xstart(
x, t, pred_noise)
46         log_variance = self.var_processor.get_variance(var_values, t) #
get the log of variance # This is not useful for DDIM, use equation
provided
47
48         ##### TODO #####
49         ##### Step 1: Implement the variance parameter 'sigma' for DDIM
sampling. #####
50         ##### Step 2: Implement x_t_minus_1 using the pred_xstart. Don't
forget #####
51         ##### to add noise for all the steps, except for the t=0.
#####
52         #####
53         ##### You may use the function 'extract_and_expand' to expand the
timestep #####
54         ##### variable 't' to the input's shape.
#####
55         ##### Assign them to the variables x_t_minus_1.
#####
56
57         ##### Start TODO #####
58         if t > 1:
59             z = torch.randn(x.shape, device=x.device)

```

```

60     else:
61         z = torch.zeros_like(x)
62
63         # alpha = extract_and_expand(self.alphas, t, x)
64         alpha_bar = extract_and_expand(self.alphas_cumprod, t, x)
65         alpha_bar_prev = extract_and_expand(self.alphas_cumprod_prev, t, x
66     )
67
68     eta = 1
69     sigma = eta * torch.sqrt((1 - alpha_bar_prev) / (1 - alpha_bar)) *
70     torch.sqrt(1 - (alpha_bar) / (alpha_bar_prev))
71
72     x_t_minus_1 = torch.sqrt(alpha_bar_prev) * pred_xstart + torch.
73     sqrt(1 - alpha_bar_prev - sigma ** 2) * pred_noise + sigma * z
74     ##### End TODO #####
75
76     return {"x_t_minus_1": x_t_minus_1, "pred_xstart": pred_xstart}
77     #####
78
79     def predict_eps_from_x_start(self, x_t, t, pred_xstart):
80         coef1 = extract_and_expand(self.sqrt_recip_alphas_cumprod, t, x_t)
81         coef2 = extract_and_expand(self.sqrt_recipm1_alphas_cumprod, t,
82         x_t)
83         return (coef1 * x_t - pred_xstart) / coef2

```

2.4.1 Result



Figure 6: DDIM Sampling Result

2.5 Task 6: Image Inpainting using RePaint

2.5.1 Code

Code submitted to Canvas

```
1 @register_sampler(name='repaint')
2 class Repaint(DDIMDiffusion):
3
4     def undo(self, image_before_step, img_after_model, est_x_0, t, debug=
5         False):
6         return self._undo(img_after_model, t)
7
8     def _undo(self, img_out, t):
9
10        beta = extract_and_expand(self.betas, t, img_out)
11
12        img_in_est = torch.sqrt(1 - beta) * img_out + \
13            torch.sqrt(beta) * torch.randn_like(img_out)
14
15        return img_in_est
16
17
18    def p_sample(
19        self,
20        model,
21        x_t_minus_one_unknown,
22        t,
23        clip_denoised=True,
24        denoised_fn=None,
25        model_kwargs=None,
26        conf=None,
27        pred_xstart=None,
28    ):
29        """
30        Sample  $x_{t-1}$  from the model at the given timestep.
31
32        :param model: the model to sample from.
33        :param x_t_minus_one_unknown: the unknown tensor at  $x_{t-1}$  (model
34        's predicted sample in the previous timestep).
35        :param t: the value of t, starting at 0 for the first diffusion
36        step.
37        :param clip_denoised: if True, clip the  $x_{start}$  prediction to [-1,
38        1].
39        :param denoised_fn: if not None, a function which applies to the
40         $x_{start}$  prediction before it is used to sample.
41        :param model_kwargs: if not None, a dict of extra keyword
42        arguments to
43        pass to the model. This can be used for conditioning.
44        :return: a dict containing the following keys:
45            - 'sample': a random sample from the model.
46            - 'pred_xstart': a prediction of  $x_0$ .
47        """
48        noise = torch.randn_like(x_t_minus_one_unknown)
```



```

45     #####
46     ##### TODO #####
47     ##### Here updated sample x_t_minus_one refers to the noisy image,
48     where the known region is #####
49     ##### obtained by adding noise to GT, and the unknown region is
50     obtained from #####
51     ##### x_t_minus_one_unknown (which is the predicted sample from
52     previous timestep) and the #####
53     ##### known and unknown region are combined using the ground
54     truth mask (gt_keep_mask). #####
55     ##### Complete the implementation to compute the updated sample (
56     x_t_minus_one) for the #####
57     ##### timestep t. Make use of the variables gt_keep_mask and gt,
58     to access the #####
59     ##### ground-truth image. and the ground-truth mask.
60     #####
61     #####
62
63     if conf["inpa_inj_sched_prev"]:
64
65         if pred_xstart is not None:
66
67             gt_keep_mask = model_kwargs['gt_keep_mask']
68             if gt_keep_mask is None:
69                 gt_keep_mask = conf.get_inpa_mask(
70 x_t_minus_one_unknown)
71
72             gt = model_kwargs['gt']
73
74             # Get x_t_minus_one_known
75             if t > 1:
76                 epsilon = torch.randn(x_t_minus_one_unknown.shape,
77 device=x_t_minus_one_unknown.device)
78             else:
79                 epsilon = torch.zeros_like(x_t_minus_one_unknown,
80 device=x_t_minus_one_unknown.device)
81             alpha_bar = extract_and_expand(self.alphas_cumprod, t,
82 x_t_minus_one_unknown)
83             x_t_minus_one_known = torch.sqrt(alpha_bar) * gt + torch.
84 sqrt(1 - alpha_bar) * epsilon
85
86             # Get x_t_minus_one
87             x_t_minus_one = gt_keep_mask * x_t_minus_one_known + (1 -
88 gt_keep_mask) * x_t_minus_one_unknown
89
90         else:
91             x_t_minus_one = x_t_minus_one_unknown
92
93     # TODO #####
94     # One-step denoising using the model: Perform a forward pass on
95     the model.

```

```

84     # Remember to scale the timestep 't' using '_scale_timesteps'
method.
85     # NOTE: Since this version of the model learns the variance along
with the score function,
86     # the output of the model would have double the number of channels
as that of the input.
87     # So assign the predicted score and variance values to the
variables below. Refer to
88     # torch.split method.
89
90     model_output = model.forward(x_t_minus_one, self._scale_timesteps(
t))
91     pred_score, var_values = torch.split(model_output, 3, dim=1)
92
93     model_mean, pred_xstart = self.mean_processor.get_mean_and_xstart(
x_t_minus_one, t, pred_score)
94     log_variance = self.var_processor.get_variance(var_values, t)    #
get the log of variance
95
96     #####
97     ##### TODO
98     #####
99     ##### Compute the noisy sample for timestep 't'
100     #####
101     ##### You should use the 'log_variance' to calculate the variance
of noise #####
102     ##### to be added.
103     #####
104     ##### Assign the sample to the variable 'sample'
105     #####
106     #####
107
108     if t > 1:
109         z = torch.randn(x_t_minus_one_unknown.shape, device=
x_t_minus_one_unknown.device)
110     else:
111         z = torch.zeros_like(x_t_minus_one_unknown, device=
x_t_minus_one_unknown.device)
112
113     sample = model_mean + torch.sqrt(torch.exp(log_variance)) * z
114
115     result = {"sample": sample,
116              "pred_xstart": pred_xstart, 'gt': model_kwargs.get('gt')}
117
118     return result
119
120 def p_sample_loop(
121     self,
122     model,
123     shape,

```

```

123         noise=None,
124         clip_denoised = True,
125         denoised_fn=None,
126         model_kwargs=None,
127         device=None,
128         progress=True,
129         return_all=False,
130         conf=None
131     ):
132         """
133         Generate samples from the model.
134
135         :param model: the model module.
136         :param shape: the shape of the samples, (N, C, H, W).
137         :param noise: if specified, the noise from the encoder to sample.
138                       Should be of the same shape as `shape`.
139         :param clip_denoised: if True, clip x_start predictions to [-1,
140 1].
141         :param denoised_fn: if not None, a function which applies to the
142                             x_start prediction before it is used to sample.
143         :param cond_fn: if not None, this is a gradient function that acts
144                         similarly to the model.
145         :param model_kwargs: if not None, a dict of extra keyword
146                             arguments to
147                             pass to the model. This can be used for conditioning.
148         :param device: if specified, the device to create the samples on.
149                       If not specified, use a model parameter's device.
150         :param progress: if True, show a tqdm progress bar.
151         :return: a non-differentiable batch of samples.
152         """
153         final = None
154         for sample in self.p_sample_loop_progressive(
155             model,
156             shape,
157             noise=noise,
158             clip_denoised=clip_denoised,
159             denoised_fn=denoised_fn,
160             model_kwargs=model_kwargs,
161             device=device,
162             progress=progress,
163             conf=conf
164         ):
165             final = sample
166
167         if return_all:
168             return final
169         else:
170             return final["sample"]
171
172     def p_sample_loop_progressive(
173         self,
174         model,
175         shape,
176         noise=None,

```

```

175         clip_denoised=True,
176         denoised_fn=None,
177         model_kwargs=None,
178         device=None,
179         progress=False,
180         conf=None
181     ):
182         """
183         Generate samples from the model and yield intermediate samples
184         from
185         each timestep of diffusion.
186
187         Arguments are the same as p_sample_loop().
188         Returns a generator over dicts, where each dict is the return
189         value of
190         p_sample().
191         """
192         if device is None:
193             device = next(model.parameters()).device
194         assert isinstance(shape, (tuple, list))
195         if noise is not None:
196             image_after_step = noise
197         else:
198             image_after_step = torch.randn(*shape, device=device)
199
200         self.gt_noises = None # reset for next image
201
202         pred_xstart = None
203
204         idx_wall = -1
205         sample_idx = defaultdict(lambda: 0)
206
207         if conf["schedule_jump_params"]:
208             times = get_schedule_jump(**conf["schedule_jump_params"])
209             time_pairs = list(zip(times[:-1], times[1:]))
210
211         if progress:
212             from tqdm.auto import tqdm
213             time_pairs = tqdm(time_pairs)
214
215         for t_last, t_cur in time_pairs:
216             idx_wall += 1
217             t_last_t = torch.tensor([t_last] * shape[0],
218                                     device=device)
219
220             if t_cur < t_last: # reverse
221                 with torch.no_grad():
222                     image_before_step = image_after_step.clone()
223                     out = self.p_sample(
224                         model,
225                         image_after_step,
226                         t_last_t,
227                         clip_denoised=clip_denoised,

```

```

227         denoised_fn=denoised_fn,
228         model_kwargs=model_kwargs,
229         conf=conf,
230         pred_xstart=pred_xstart
231     )
232     image_after_step = out["sample"]
233     pred_xstart = out["pred_xstart"]
234
235     sample_idxes[t_cur] += 1
236
237     yield out
238
239 else:
240     t_shift = conf.get('inpa_inj_time_shift', 1)
241
242     image_before_step = image_after_step.clone()
243     image_after_step = self.undo(
244         image_before_step, image_after_step,
245         est_x_0=out['pred_xstart'], t=t_last_t+t_shift,
debug=False)
246     pred_xstart = out["pred_xstart"]

```

2.5.2 Result



Figure 7: RePaint Inpainting Result

2.6 Task 7: Image Inpainting using Diffusion Posterior Sampling

2.6.1 Code

Code submitted to Canvas.

```

1 @register_conditioning_method(name='ps')
2 class PosteriorSampling(ConditioningMethod):
3     def __init__(self, operator, noiser, **kwargs):
4         super().__init__(operator, noiser)
5         self.scale = kwargs.get('scale', 1.0)
6
7     def conditioning(self, x_i, x_t_minus_one, x_0_hat, measurement, **
8 kwargs):
9         """
10         The conditioning function as shown in line 7
11
12         Args:
13             x_i: torch.Tensor, x_i
14             x_t_minus_one, torch.Tensor, x_t_minus_1 prime
15             x_0_hat: torch.Tensor, predicted x_0
16             measurement: torch.Tensor, y, the corrupted image
17
18         """
19         # norm_grad, norm = self.grad_and_value(x_prev=x_prev, x_0_hat=
20 x_0_hat, measurement=measurement, **kwargs)
21         ##### Start TODO #####
22         ##### Implement the conditional sampling in line 7 #####
23         ##### A(x_0_hat) is already provided to you as A #####
24         ##### Also torch.autograd.grad() is provided to you to calculate
25 the gradient of the
26         ##### norm term with respect to x_i, you can check https://
27 pytorch.org/docs/stable/generated/torch.autograd.grad.html#torch.
28 autograd.grad
29         ##### for its detailed usage. You only need to specify the
30 outputs and inputs here.
31         A = self.operator.forward(x_0_hat, **kwargs)
32         new_x_t_minus_one = None
33         difference = None
34         norm = None
35         diff_output = None # outputs of the differentiated function
36         diff_input = None # Inputs w.r.t. which the gradient will be
37 returned
38
39         # My code
40         difference = measurement - A
41         norm = torch.norm(difference)
42         diff_output = norm
43         diff_input = x_i
44
45         ## TODO: Don't delete this line, you will use this
46         norm_grad = torch.autograd.grad(outputs=diff_output, inputs=
47 diff_input)[0]
48
49         new_x_t_minus_one = x_t_minus_one - self.scale * norm_grad
50
51         ##### END TODO #####
52         return new_x_t_minus_one

```

2.6.2 Result



Figure 8: DPS Inpainting Result

3 Appendix

Full Notebook pdf given in next page

Submitted by Wensong Hu on April 4th, 2024.