

April 4th, 2024

HW5 — Generative Models

1 Section 1: Pix2Pix

1.1 Task 1: Dataloading

Code submitted to Canvas. Full notebook see appendix.

```
1 class Edges2Image(Dataset):
2     def __init__(self, root_dir, split='train', transform=None):
3         """
4             Args:
5                 root_dir: the directory of the dataset
6                 split: "train" or "val"
7                 transform: pytorch transformations.
8         """
9
10        self.transform = transform
11
12        self.files = glob.glob(os.path.join(root_dir, split, '*.jpg'))
13
14    def __len__(self):
15        return len(self.files)
16
17    def __getitem__(self, idx):
18        img = Image.open(self.files[idx])
19        img = np.asarray(img)
20        if self.transform:
21            img = self.transform(img)
22        return img
23
24    transform = transforms.Compose([
25        transforms.ToTensor(),
26        transforms.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
27    ])
28
29    #
30    ##### TODO: Construct the dataloader
31    #
32    # For the train_loader, please use a batch size of 4 and set shuffle True
33    #
```

```

32 # For the val_loader, please use a batch size of 5 and set shuffle False
# Hint: You'll need to create instances of the class above, name them as
# tr_dt and te_dt. The dataloaders should be named as train_loader and
# test_loader. You also need to include transform in your class
# instances
#
#
37 #####
38
39 tr_dt = Edges2Image(root_dir="mini-edges2shoes", split='train', transform=
    transform)
40 te_dt = Edges2Image(root_dir="mini-edges2shoes", split='val', transform=
    transform)
41
42 train_loader = DataLoader(tr_dt, batch_size=4, shuffle=True)
43 test_loader = DataLoader(te_dt, batch_size=5, shuffle=True)
44
45 #####
46 #                         END OF YOUR CODE
47 #
48 #####
49 # Make sure that you have 1,000 training images and 100 testing images
# before moving on
50 print('Number of training images {}, number of testing images {}'.format(
    len(tr_dt), len(te_dt)))

```

1.2 Task 2: Training Pix2Pix

1.2.1 Codes

Code submitted to Canvas. Full notebook see appendix.

```

1 # Hint: you could use following loss to complete following function
2 BCE_loss = nn.BCELoss().cuda()
3 L1_loss = nn.L1Loss().cuda()
4
5 def train(G, D, num_epochs = 20):
6     hist_D_losses = []
7     hist_G_losses = []
8     hist_G_L1_losses = []
9     #
10 #####

```

```

10 # TODO: Add Adam optimizer to generator and discriminator
11 #
12 # You will use lr=0.0002, beta=0.5, beta2=0.999
13 #
14 #####
15
16
17 #
18 # END OF YOUR CODE
19 #
20 #####
21 print('training start!')
22 start_time = time.time()
23 for epoch in range(num_epochs):
24     print('Start training epoch %d' % (epoch + 1))
25     D_losses = []
26     G_losses = []
27     epoch_start_time = time.time()
28     num_iter = 0
29     for x_ in train_loader:
30         y_ = x_[:, :, :, img_size:]
31         x_ = x_[:, :, :, 0:img_size]
32
33         x_, y_ = x_.cuda(), y_.cuda()
34
35         #####
36         # TODO: Implement training code for the discriminator.
37         #
38         # Recall that the loss is the mean of the loss for real images and
39         # fake
40         # images, and made by some calculations with zeros and ones
41         #
42         # We have defined the BCE_loss, which you might would like to use.
43         #
44         #
45         # NOTE: While training the Discriminator, the output of the
46         # generator
47         #
48         # must be detached from the computational graph. Refer to the method
49         #
50         # torch.Tensor.detach()
51         #
52         #
53 #####

```

```

44
45     N = x_.shape[0]
46     # Generate data
47     fake_data = G.forward(x_).detach()
48
49     #1. Train the discriminator
50     # D real data BCE loss
51     D_real_preds = D.forward(torch.cat((x_, y_), dim=1))
52     D_y_real = torch.ones_like(D_real_preds)
53     # D_real_loss = torch.sum(torch.log(D_real_preds))
54     D_real_loss = BCE_loss(D_real_preds, D_y_real)
55
56     # D fake data BCE loss
57     D_fake_preds = D.forward(torch.cat((x_, fake_data), dim=1))
58     D_y_fake = torch.zeros_like(D_fake_preds)
59     # D_fake_loss = torch.sum(torch.log(1 - D_fake_preds))
60     D_fake_loss = BCE_loss(D_fake_preds, D_y_fake)
61
62     # D loss
63     loss_D = D_real_loss + D_fake_loss
64
65     # Train D
66     D_optimizer.zero_grad()
67     loss_D.backward()
68     D_optimizer.step()
69
70     #
71     ##### END OF YOUR CODE #####
72     #
73     #
74     ##### TODO: Implement training code for the Generator.
75     #
76     #
77     # 1. Train the generator
78     # 2. Append the losses to the lists 'hist_G_L1_losses' and '
79     hist_G_losses'
80     # (Only append the data to the list, not the complete tensor, refer
81     # torch.Tensor.item()).
82
83     # Generate data
84     fake_data = G.forward(x_)
85

```

```

86     # 1. Train the generator
87     # G BCE loss
88     G_fake_preds = D.forward(torch.cat((x_, fake_data), dim=1))
89     G_y_fake = torch.zeros_like(G_fake_preds)
90     G_bce_loss = BCE_loss(G_fake_preds, G_y_fake)
91
92     # G l1 loss
93     G_l1_loss = L1_loss(fake_data, y_)
94
95     # G loss
96     lamb = 100
97     loss_G = G_bce_loss + lamb * G_l1_loss
98
99     # Train G
100    G_optimizer.zero_grad()
101    loss_G.backward()
102    G_optimizer.step()
103
104    # 2. Append the losses to the lists 'hist_G_L1_losses' and '
105    hist_D_losses'
106    # (Only append the data to the list, not the complete tensor, refer
107    # to torch.Tensor.item()).
108    hist_G_losses.append(G_bce_loss.detach().item())
109    hist_G_L1_losses.append(G_l1_loss.detach().item())
110    #
111    ######
112
113    D_losses.append(loss_D.detach().item())
114    hist_D_losses.append(loss_D.detach().item())
115    G_losses.append(loss_G)
116    num_iter += 1
117
118
119
120    epoch_end_time = time.time()
121    per_epoch_ptime = epoch_end_time - epoch_start_time
122
123    print('[%d/%d] - using time: %.2f seconds' % ((epoch + 1), num_epochs,
124    per_epoch_ptime))
124    print('loss of discriminator D: %.3f' % (torch.mean(torch.FloatTensor(
125    D_losses))))
125    print('loss of generator G: %.3f' % (torch.mean(torch.FloatTensor(
126    G_losses))))
126    if epoch == 0 or (epoch + 1) % 5 == 0:
127        with torch.no_grad():
128            show_result(G, fixed_x_, fixed_y_, (epoch+1))
129
130    end_time = time.time()

```

```
131     total_ptime = end_time - start_time  
132  
133     return hist_D_losses, hist_G_losses, hist_G_L1_losses
```

1.2.2 Report Results



Figure 1: Result Visualization after 20 training epoches

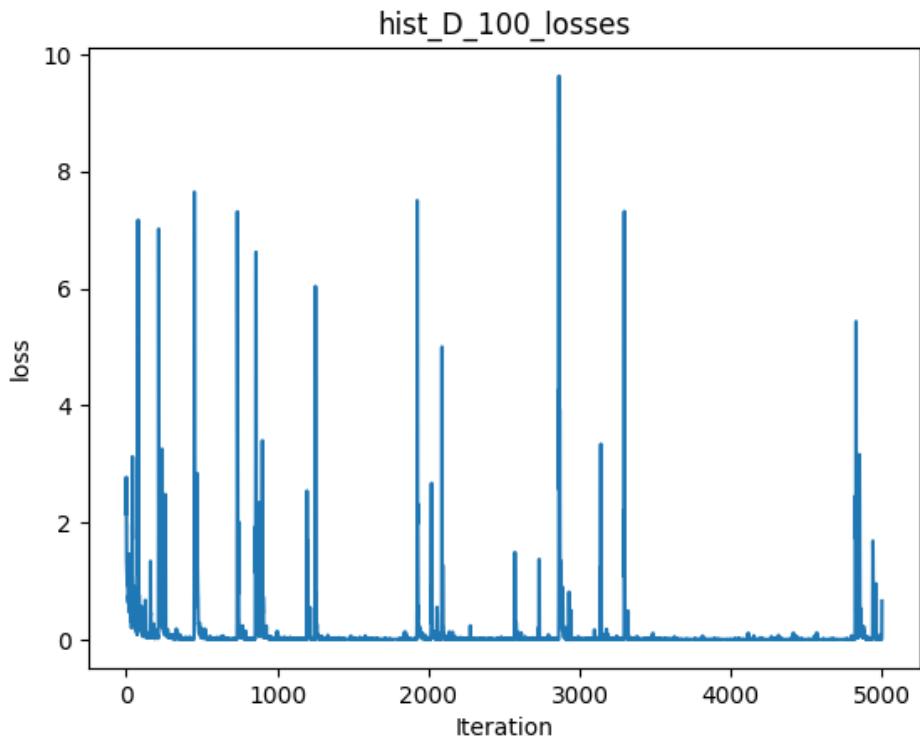


Figure 2: Discriminator BCE loss

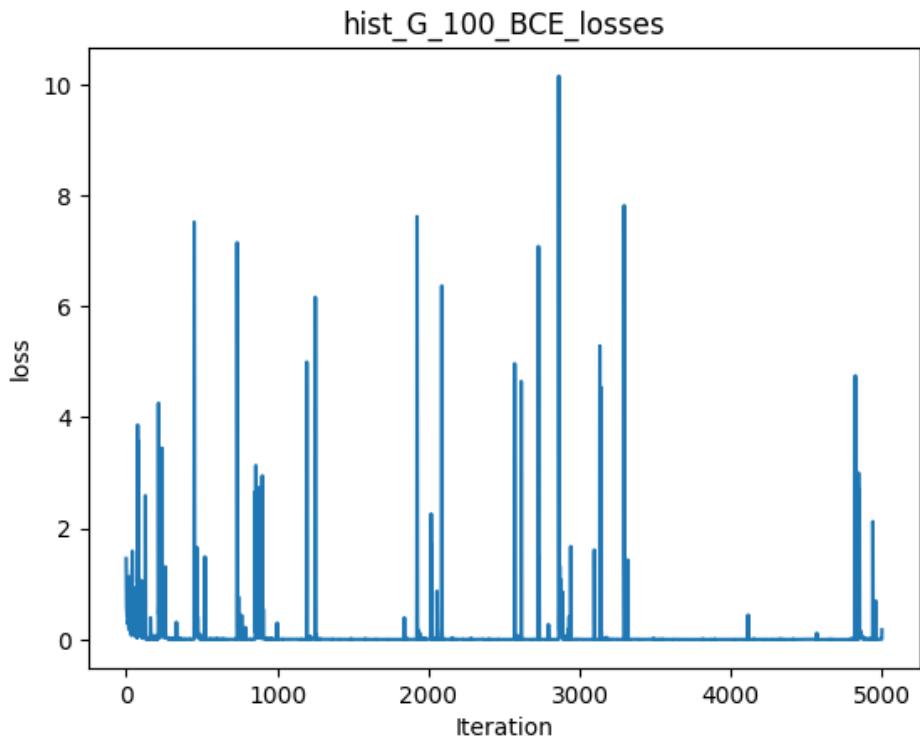


Figure 3: Generator BCE loss

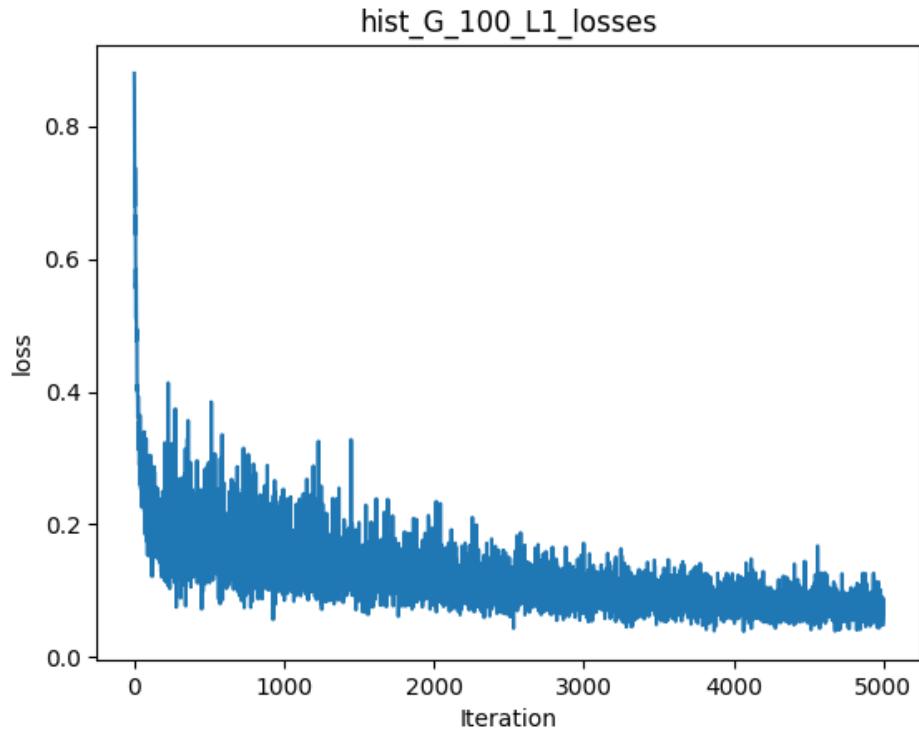


Figure 4: Generator L1 loss

2 Section 2: Diffusion Models

2.1 Task 3: Unconditional Sampling using DDPM: get_name_beta_schedule

Code submitted to Canvas.

```
1 def get_named_beta_schedule(schedule_name, num_diffusion_timesteps,
2     beta_min=0.0001, beta_max=0.02):
3     """
4         Get a pre-defined beta schedule for the given name.
5
6     Args:
7         schedule_name: str, name of the variance schedule, 'linear' or 'cosine'
8         num_diffusion_timesteps: int, number of the entire diffusion
9             timesteps
10        beta_min: float, minimum value of beta
11        beta_max: float, maximum value of beta
12
13    Returns:
14        betas: np.ndarray, a 1-d array of size num_diffusion_timesteps,
15        contains all the beta for each timestep
16
17    """
18    betas = None
19    if schedule_name == "linear":
20        ##### START TODO #####
21
```

```

18     # Implement the linear schedule
19     # Uniformly divide the [beta_min, beta_max) to
20     num_diffusion_timesteps values.
21     betas = np.linspace(beta_min, beta_max, num_diffusion_timesteps)
22     ##### END TODO #####
23
24 elif schedule_name == "cosine":
25     ##### START TODO #####
26     # Implement the cosine schedule
27     # Assume s = 0.008 and beta_clip=0.999
28     s = 0.008
29     beta_clip = 0.999
30
31     betas = np.zeros(num_diffusion_timesteps)
32     f_0 = 0.0
33     alpha_bar_tminus1 = 1.0
34     for t in range(num_diffusion_timesteps):
35         if t == 0:
36             f_0 = (np.cos((0/num_diffusion_timesteps + s) / (1 + s) *
37 np.pi / 2)) ** 2
38             # alpha_bar_tminus1 = 1
39
40         else:
41             f_t = (np.cos((t/num_diffusion_timesteps + s) / (1 + s) *
42 np.pi / 2)) ** 2
43             alpha_bar_t = f_t / f_0
44             # import pdb; pdb.set_trace()
45             betas[t-1] = np.clip((1 - alpha_bar_t / alpha_bar_tminus1)
46 , a_min=0, a_max=beta_clip)
47             alpha_bar_tminus1 = alpha_bar_t
48             # import pdb; pdb.set_trace()
49             betas[-1] = beta_clip
50
51     ##### End TODO #####
52 else:
53     raise NotImplementedError(f"unknown beta schedule: {schedule_name}")
54
55 return betas

```

2.2 Task 4: Unconditional Sampling using DDPM: DDPM

2.2.1 Code

Code submitted to Canvas:

```

1 @register_sampler("ddpm")
2 class DDPMdiffusion:
3     def __init__(self,
4                  betas,
5                  dynamic_threshold,
6                  clip_denoised,
7                  rescale_timesteps,

```

```

8         **kwargs
9     ):
10
11     # use float64 for accuracy.
12     betas = np.array(betas, dtype=np.float64)
13     self.betas = betas
14     assert self.betas.ndim == 1, "betas must be 1-D"
15     assert (0 < self.betas).all() and (self.betas <=1).all(), "betas
16     must be in (0..1]"
17
18     self.num_timesteps = int(self.betas.shape[0])
19     self.rescale_timesteps = rescale_timesteps
20
21     ##### START TODO #####
22     # Calculate the values of alpha
23     # Also we will need the cumulated product of alpha.
24     # And during sampling we need the value of cumulated product of
25     # alpha from
26     # previous or next timestep.
27     self.alphas = 1 - betas
28     self.alphas_cumprod = np.cumprod(self.alphas) # cumulated
29     product of alphas
30     self.alphas_cumprod_prev = np.concatenate((np.array([1]), self.
31     alphas_cumprod[:-1])) # first T-1 elements of alphas_cumprod, append
32     1.0 at the begining, to make its length T
33     self.alphas_cumprod_next = np.concatenate((self.alphas_cumprod
34     [1:], np.array([0]))) # last T-1 elements of alphas_cumprod, append 0.0
35     at the end, to make its length T
36
37     ##### END TODO #####
38     assert self.alphas_cumprod_prev.shape == (self.num_timesteps,)
39
40     self.mean_processor = EpsilonXMeanProcessor(betas=betas,
41                                                 dynamic_threshold=
42     dynamic_threshold,
43                                                 clip_denoised=
44     clip_denoised)
45
46     self.var_processor = LearnedRangeVarianceProcessor(betas=betas)
47
48     def p_sample_loop(self,
49                         model,
50                         x_start,
51                         record,
52                         save_root,
53                         measurement=None,
54                         measurement_cond_fn=None,
55                         uncond=False):
56         """
57             The function used for sampling from noise.
58
59         Args:
60             model: nn.Module, the pretrained model that is used to predict
61             the score and variance

```

```

52         x_start: torch.Tensor, random noise input
53         measurement: torch.Tensor, our corrupted observation
54         measurement_cond_fn: conditional function used to perform
55         conditional sampling, is None for unconditional sampling
56         record: Bool, save intermediate results if True
57         save_root: str, root of the directory to save the results
58         uncond: Bool, perform unconditional sampling if True, else
59         perform conditional sampling
60         """
61     if not uncond:
62         assert measurement is not None and measurement_cond_fn is not
63     None, \
64         "measurement and measurement conditional function is
65     required for conditional sampling"
66
67     img = x_start    # start from random noise
68     device = x_start.device
69
70     ##### Start TODO #####
71     # Implement the sample loop
72     # Call p_sample for every iteration
73     # It requires only one line of code implementation here
74
75     pbar = tqdm(list(range(self.num_timesteps))[:-1])
76     for idx in pbar:
77         time = torch.tensor([idx] * img.shape[0], device=device)
78
79         img = self.p_sample(model=model, x=img, t=time) ['x_t_minus_1']
80
81         img = img.detach_()
82
83         if record:
84             if idx % 10 == 0:
85                 file_path = os.path.join(save_root, f"progress/x_{str(
86                     idx).zfill(4)}.png")
87                 plt.imsave(file_path, clear_color(img))
88
89     return img
90
91 def p_sample(self, model, x, t):
92     """
93         Posterior sampling process, when given the model, x_t and timestep
94         t, it returns predicted
95         x_0 and x_t_minus_1
96
97         We have already provided you with the function to get the log of
98         the variance.
99         Use self.var_processor.get_variance(var_values, t), where
100         var_values is
101             the 3:6 channels of the direct output of the model.
102             example usage: log_variance = self.var_processor.get_variance(
103                 var_values, t)
104
105

```

```

96     You can also use the helper function extract_and_expand() to
97     extract the value
98     corresponding to timestep and expand it to the same size as the
99     target for broadcast.
100    example usage: coef1 = extract_and_expand(self.
101 posterior_mean_coef1, t, x_start)
102
103    Args:
104        model: nn.Module, the UNet model, you can call model(x, t) to
105        get the output tensor with size (B, 6, H, W)
106        x: torch.Tensor, shape (1, 3, H, W), x_t
107        t: torch.Tenosr, shape (1,), timestep
108
109    Returns:
110        output_dict: dict, contains predicted x_t_minus_1 and x_0
111        """
112    ##### Start TODO#####
113    ##### Get the predicted score and variance of the pretrained model
114    #####
115    model_output = model.forward(x, t)
116    pred_noise = model_output[:, :3, :, :]
117    var_values = model_output[:, 3:, :, :]
118    ##### End TODO #####
119
120    log_variance = self.var_processor.get_variance(var_values, t)    #
121    get the log of variance
122
123    ##### Start TODO #####
124    ##### get predicted x_0 and x_t_minus_1 #####
125    ##### don't forget to add noise for all the steps, except for the
126    last one #####
127    if t > 1:
128        z = torch.randn(x.shape, dtype=x.dtype, device=x.device)
129    else:
130        z = torch.zeros_like(x, device=x.device)
131    # import pdb; pdb.set_trace()
132    alpha = extract_and_expand(self.alphas, t, x)
133    alpha_bar = extract_and_expand(self.alphas_cumprod, t, x)
134    x_t_minus_1 = (1 / torch.sqrt(alpha)) * (x - ((1 - alpha) / torch.
sqrt(1 - alpha_bar)) * pred_noise) + torch.sqrt(torch.exp(log_variance)
) * z
135
136    ##### End TODO #####
137
138    assert x_t_minus_1.shape == log_variance.shape == x.shape
139
140    output_dict = {'x_t_minus_1': x_t_minus_1}
141    return output_dict

```

2.2.2 Result



Figure 5: DDPM Sampling Result

2.3 Task 5: Unconditional Sampling using DDIM

2.4 Code

Code submitted to Canvas.

```
1 @register_sampler("ddim")
2 class DDIMDiffusion(DDPMDiffusion):
3
4     def __init__(self, use_timesteps, **kwargs):
5
6         self.timestep_map = []
7         self.original_num_steps = len(kwargs["betas"])
8
9         base_alphas_cumprod = DDPMDiffusion(**kwargs).alphas_cumprod # pylint: disable=missing-kwoa
10        last_alpha_cumprod = 1.0
11        new_betas = []
12        self.use_timesteps = set(use_timesteps)
13
14        for i, alpha_cumprod in enumerate(base_alphas_cumprod):
15            if i in self.use_timesteps:
16                new_betas.append(1 - alpha_cumprod / last_alpha_cumprod)
17                last_alpha_cumprod = alpha_cumprod
18                self.timestep_map.append(i)
19        kwargs["betas"] = np.array(new_betas)
20        super().__init__(**kwargs)
21
22
```

```

23     def _scale_timesteps(self, t):
24         if self.rescale_timesteps:
25             return t.float() * (self.original_num_steps / self.
26         num_timesteps)
27         return t
28
29     def p_sample(self, model, x, t, eta=0.0):
30         ##### TODO #####
31         ##### Get the predicted score and variance of the pretrained model
32         ##### Don't forget to use _scale_timesteps to scale the timestep
33         for calling the model prediction.
34         ##### You don't need to scale the timestep for further
35         computations of x_t_minus_1.
36         ##### NOTE: Since this version of the model learns the variance
37         along with the score function,
38         ##### the output of the model would have double the number of
39         channels as that of the input.
40         ##### So assign the predicted score and variance values to the
41         variables below. Refer to
42         ##### torch.split method.
43         ##### Start TODO #####
44         model_output = model.forward(x, self._scale_timesteps(t))
45         # import pdb; pdb.set_trace()
46         pred_noise, var_values = torch.split(model_output, 3, dim=1)
47         ##### End TODO #####
48
49         model_mean, pred_xstart = self.mean_processor.get_mean_and_xstart(
50             x, t, pred_noise)
51         log_variance = self.var_processor.get_variance(var_values, t)    #
52         get the log of variance # This is not useful for DDIM, use equation
53         provided
54
55         ##### TODO #####
56         ##### Step 1: Implement the variance parameter 'sigma' for DDIM
57         sampling. #####
58         ##### Step 2: Implement x_t_minus_1 using the pred_xstart. Don't
59         forget #####
60         ##### to add noise for all the steps, except for the t=0.
61         #####
62         #####
63         ##### You may use the function 'extract_and_expand' to expand the
64         timestep #####
65         ##### variable 't' to the input's shape.
66         #####
67         ##### Assign them to the variables x_t_minus_1.
68         #####
69
70         ##### Start TODO #####
71         if t > 1:
72             z = torch.randn(x.shape, device=x.device)

```

```

60     else:
61         z = torch.zeros_like(x)
62
63     # alpha = extract_and_expand(self.alphas, t, x)
64     alpha_bar = extract_and_expand(self.alphas_cumprod, t, x)
65     alpha_bar_prev = extract_and_expand(self.alphas_cumprod_prev, t, x
66 )
67
68     eta = 1
69     sigma = eta * torch.sqrt((1 - alpha_bar_prev) / (1 - alpha_bar)) *
70     torch.sqrt(1 - (alpha_bar) / (alpha_bar_prev))
71     ##### End TODO #####
72
73     return {"x_t_minus_1": x_t_minus_1, "pred_xstart": pred_xstart}
74 ##########
75
76
77     def predict_eps_from_x_start(self, x_t, t, pred_xstart):
78         coef1 = extract_and_expand(self.sqrt_recip_alphas_cumprod, t, x_t)
79         coef2 = extract_and_expand(self.sqrt_recipm1_alphas_cumprod, t,
80 x_t)
81         return (coef1 * x_t - pred_xstart) / coef2

```

2.4.1 Result



Figure 6: DDIM Sampling Result

2.5 Task 6: Image Inpainting using RePaint

2.5.1 Code

Code submitted to Canvas

```
1 @register_sampler(name='repaint')
2 class Repaint(DDIMDiffusion):
3
4     def undo(self, image_before_step, img_after_model, est_x_0, t, debug=False):
5         return self._undo(img_after_model, t)
6
7
8     def _undo(self, img_out, t):
9
10        beta = extract_and_expand(self.betas, t, img_out)
11
12        img_in_est = torch.sqrt(1 - beta) * img_out + \
13            torch.sqrt(beta) * torch.randn_like(img_out)
14
15        return img_in_est
16
17
18    def p_sample(
19        self,
20        model,
21        x_t_minus_one_unknown,
22        t,
23        clip_denoised=True,
24        denoised_fn=None,
25        model_kwargs=None,
26        conf=None,
27        pred_xstart=None,
28    ):
29        """
30            Sample x_{t-1} from the model at the given timestep.
31
32            :param model: the model to sample from.
33            :param x_t_minus_one_unknown: the unknown tensor at x_{t-1} (model
34            's predicted sample in the previous timestep).
35            :param t: the value of t, starting at 0 for the first diffusion
36            step.
37            :param clip_denoised: if True, clip the x_start prediction to [-1,
38            1].
39            :param denoised_fn: if not None, a function which applies to the
40            x_start prediction before it is used to sample.
41            :param model_kwargs: if not None, a dict of extra keyword
42            arguments to
43            pass to the model. This can be used for conditioning.
44            :return: a dict containing the following keys:
45                    - 'sample': a random sample from the model.
46                    - 'pred_xstart': a prediction of x_0.
47        """
48        noise = torch.randn_like(x_t_minus_one_unknown)
```

```

45
46     ##### TODO #####
47     ##### Here updated sample x_t_minus_one refers to the noisy image ,
48     where the known region is #####
49     ##### obtained by adding noise to GT, and the unknown region is
50     ##### obtained from #####
51     ##### x_t_minus_one_unknown (which is the predicted sample from
52     previous timestep) and the #####
53     ##### known and unknown region are combined using the ground
54     truth mask (gt_keep_mask). #####
55     ##### Compelete the implementation to compute the updated sample (
56     x_t_minus_one) for the #####
57     ##### timestep t. Make use of the variables gt_keep_mask and gt,
58     to access the #####
59     ##### ground-truth image. and the ground-truth mask.
60     #####
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83

```

```

84     # Remember to scale the timestep 't' using '_scale_timesteps'
85     # method.
86     # NOTE: Since this version of the model learns the variance along
87     # with the score function,
88     # the output of the model would have double the number of channels
89     # as that of the input.
90     # So assign the predicted score and variance values to the
91     # variables below. Refer to
92     # torch.split method.
93
94     model_output = model.forward(x_t_minus_one, self._scale_timesteps(
95         t))
96     pred_score, var_values = torch.split(model_output, 3, dim=1)
97
98     model_mean, pred_xstart = self.mean_processor.get_mean_and_xstart(
99         x_t_minus_one, t, pred_score)
100    log_variance = self.var_processor.get_variance(var_values, t)      # get the log of variance
101
102    ##### TODO
103    #####
104    ##### Compute the noisy sample for timestep 't'
105    #####
106    ##### You should use the 'log_variance' to calculate the variance
107    # of noise #####
108    ##### to be added.
109    #####
110    ##### Assign the sample to the variable 'sample'
111    #####
112    ##### #####
113    result = {"sample": sample,
114              "pred_xstart": pred_xstart, 'gt': model_kwargs.get('gt')}
115
116    return result
117
118
119    def p_sample_loop(
120        self,
121        model,
122        shape,

```

```

123     noise=None ,
124     clip_denoised = True ,
125     denoised_fn=None ,
126     model_kwargs=None ,
127     device=None ,
128     progress=True ,
129     return_all=False ,
130     conf=None
131 ) :
132 """
133 Generate samples from the model.
134
135 :param model: the model module.
136 :param shape: the shape of the samples, (N, C, H, W).
137 :param noise: if specified, the noise from the encoder to sample.
138             Should be of the same shape as `shape`.
139 :param clip_denoised: if True, clip x_start predictions to [-1,
140 ].           1].
141 :param denoised_fn: if not None, a function which applies to the
142             x_start prediction before it is used to sample.
143 :param cond_fn: if not None, this is a gradient function that acts
144             similarly to the model.
145 :param model_kwargs: if not None, a dict of extra keyword
146 arguments to
147             pass to the model. This can be used for conditioning.
148 :param device: if specified, the device to create the samples on.
149             If not specified, use a model parameter's device.
150 :param progress: if True, show a tqdm progress bar.
151 :return: a non-differentiable batch of samples.
152 """
153     final = None
154     for sample in self.p_sample_loop_progressive(
155         model,
156         shape,
157         noise=noise,
158         clip_denoised=clip_denoised,
159         denoised_fn=denoised_fn,
160         model_kwargs=model_kwargs,
161         device=device,
162         progress=progress,
163         conf=conf
164     ):
165         final = sample
166
167     if return_all:
168         return final
169     else:
170         return final["sample"]
171
172 def p_sample_loop_progressive(
173     self,
174     model,
175     shape,
176     noise=None ,

```

```

175     clip_denoised=True ,
176     denoised_fn=None ,
177     model_kwargs=None ,
178     device=None ,
179     progress=False ,
180     conf=None
181 ) :
182     """
183     Generate samples from the model and yield intermediate samples
184     from
185     each timestep of diffusion.
186
187     Arguments are the same as p_sample_loop().
188     Returns a generator over dicts, where each dict is the return
189     value of
190     p_sample().
191     """
192
193     if device is None:
194         device = next(model.parameters()).device
195     assert isinstance(shape, (tuple, list))
196     if noise is not None:
197         image_after_step = noise
198     else:
199         image_after_step = torch.randn(*shape, device=device)
200
201     self.gt_noises = None # reset for next image
202
203     pred_xstart = None
204
205     idx_wall = -1
206     sample_idxs = defaultdict(lambda: 0)
207
208     if conf["schedule_jump_params"]:
209         times = get_schedule_jump(**conf["schedule_jump_params"])
210         time_pairs = list(zip(times[:-1], times[1:]))
211
212         if progress:
213             from tqdm.auto import tqdm
214             time_pairs = tqdm(time_pairs)
215
216         for t_last, t_cur in time_pairs:
217             idx_wall += 1
218             t_last_t = torch.tensor([t_last] * shape[0],
219                                    device=device)
220
221             if t_cur < t_last: # reverse
222                 with torch.no_grad():
223                     image_before_step = image_after_step.clone()
224                     out = self.p_sample(
225                         model,
226                         image_after_step,
227                         t_last_t,
228                         clip_denoised=clip_denoised,

```

```

227             denoised_fn=denoised_fn ,
228             model_kwargs=model_kwargs ,
229             conf=conf ,
230             pred_xstart=pred_xstart
231         )
232         image_after_step = out["sample"]
233         pred_xstart = out["pred_xstart"]
234
235         sample_idxs[t_cur] += 1
236
237         yield out
238
239     else:
240         t_shift = conf.get('inpa_inj_time_shift', 1)
241
242         image_before_step = image_after_step.clone()
243         image_after_step = self.undo(
244             image_before_step, image_after_step,
245             est_x_0=out['pred_xstart'], t=t_last_t+t_shift,
246             debug=False)
246             pred_xstart = out["pred_xstart"]

```

2.5.2 Result



Figure 7: RePaint Inpainting Result

2.6 Task 7: Image Inpainting using Diffusion Posterior Sampling

2.6.1 Code

Code submitted to Canvas.

```

1 @register_conditioning_method(name='ps')
2 class PosteriorSampling(ConditioningMethod):
3     def __init__(self, operator, noiser, **kwargs):
4         super().__init__(operator, noiser)
5         self.scale = kwargs.get('scale', 1.0)
6
7     def conditioning(self, x_i, x_t_minus_one, x_0_hat, measurement, **kwargs):
8         """
9             The conditioning function as shown in line 7
10
11         Args:
12             x_i: torch.Tensor, x_i
13             x_t_minus_one, torch.Tensor, x_t_minus_1 prime
14             x_0_hat: torch.Tensor, predicted x_0
15             measurement: torch.Tensor, y, the corrupted image
16         """
17         # norm_grad, norm = self.grad_and_value(x_prev=x_prev, x_0_hat=
18         # x_0_hat, measurement=measurement, **kwargs)
19         ##### Start TODO #####
20         ##### Implement the conditional sampling in line 7 #####
21         ##### A(x_0_hat) is already provided to you as A #####
22         ##### Also torch.autograd.grad() is provided to you to calculate
the gredient of the
23         ##### norm term with respect to x_i, you can check https://
pytorch.org/docs/stable/generated/torch.autograd.grad.html#torch.
autograd.grad
24         ##### for its detailed usage. You only need to specify the
outputs and inputs here.
25         A = self.operator.forward(x_0_hat, **kwargs)
26         new_x_t_minus_one = None
27         difference = None
28         norm = None
29         diff_output = None # outputs of the differentiated function
30         diff_input = None # Inputs w.r.t. which the gradient will be
returned
31
32         # My code
33         difference = measurement - A
34         norm = torch.norm(difference)
35         diff_output = norm
36         diff_input = x_i
37
38         ## TODO: Don't delete this line, you will use this
39         norm_grad = torch.autograd.grad(outputs=diff_output, inputs=
diff_input)[0]
40
41         new_x_t_minus_one = x_t_minus_one - self.scale * norm_grad
42
43         ##### END TODO #####
44         return new_x_t_minus_one

```

2.6.2 Result



Figure 8: DPS Inpainting Result

3 Appendix

Full Notebook pdf given in next page

EECS 442 PS5: Generative Adversarial Models

Please provide the following information (e.g. Jinfan Zhou, zjf): Wensong Hu, umhws

Starting

Run the following code to import the modules you'll need. After you finish the assignment, remember to run all cells and convert the notebook to a .pdf file for Gradescope submission.

```
In [1]: !pip install torchsummary
import pickle
import numpy as np
import matplotlib.pyplot as plt
import os
import time
import itertools
from matplotlib import image
import glob as glob
from PIL import Image

import torch
import torchvision
from torchvision import datasets, models, transforms
import torch.nn as nn
import torch.optim as optim
from torch.autograd import Variable
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
from torchsummary import summary

print("PyTorch Version: ",torch.__version__)
print("Torchvision Version: ",torchvision.__version__)
# Detect if we have a GPU available
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
if torch.cuda.is_available():
    print("Using the GPU!")
else:
    print("WARNING: Could not find GPU! Using CPU only. If you want to enable
```

```
Requirement already satisfied: torchsummary in /home/umhws/anaconda3/envs/ee
cs442/lib/python3.10/site-packages (1.5.1)
PyTorch Version: 2.2.1+cu118
Torchvision Version: 0.17.1+cu118
Using the GPU!
```

Problem 6.1 pix2pix

You will build pix2pix for image translation.

In this question, you will need to:

1. Construct dataloaders for train/test datasets
2. Build Generator and Discriminator
3. Train pix2pix and visualize the results during training
4. Plot the loss of generator/discriminator v.s. iteration
5. Design your own shoes (optional)

Step 0: Downloading the dataset.

We first download the `mini-edges2shoes` dataset sampled from the original `edges2shoes` dataset. The `mini-edges2shoes` dataset contains 1,000 training image pairs, and 100 testing image pairs.

There's nothing you need to implement for this part.

```
In [ ]: # Download the mini-edges2shoes dataset
!rm -r mini-edges2shoes.zip
!rm -r mini-edges2shoes
!wget http://www.eecs.umich.edu/courses/eecs442-ahowens/mini-edges2shoes.zip
!unzip -q mini-edges2shoes.zip

!nvidia-smi
```

Step 1: Build dataloaders for train and test

We will first build dataloaders with PyTorch built-in classes. Here we build a Custom Dataset rather than a built-in dataset.

```
In [3]: class Edges2Image(Dataset):
    def __init__(self, root_dir, split='train', transform=None):
        """
        Args:
            root_dir: the directory of the dataset
            split: "train" or "val"
            transform: pytorch transformations.
        """

        self.transform = transform
```

```

    self.files = glob.glob(os.path.join(root_dir, split, '*.jpg'))

    def __len__(self):
        return len(self.files)

    def __getitem__(self, idx):
        img = Image.open(self.files[idx])
        img = np.asarray(img)
        if self.transform:
            img = self.transform(img)
        return img

    transform = transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
    ])

#####
# TODO: Construct the dataloader
# For the train_loader, please use a batch size of 4 and set shuffle True
# For the val_loader, please use a batch size of 5 and set shuffle False
# Hint: You'll need to create instances of the class above, name them as
# tr_dt and te_dt. The dataloaders should be named as train_loader and
# test_loader. You also need to include transform in your class
# instances
#####

tr_dt = Edges2Image(root_dir="mini-edges2shoes", split='train', transform=transform)
te_dt = Edges2Image(root_dir="mini-edges2shoes", split='val', transform=transform)

train_loader = DataLoader(tr_dt, batch_size=4, shuffle=True)
test_loader = DataLoader(te_dt, batch_size=5, shuffle=True)

#####
# END OF YOUR CODE
#####

# Make sure that you have 1,000 training images and 100 testing images before
print('Number of training images {}, number of testing images {}'.format(len

```

Number of training images 1000, number of testing images 100

In [4]:

```

#Sample Output used for visualization
test = test_loader.__iter__().__next__()
img_size = 256
fixed_y_ = test[:, :, :, img_size: ].cuda()
fixed_x_ = test[:, :, :, 0:img_size].cuda()
print(len(train_loader))
print(len(test_loader))
print(fixed_y_.shape)

# plot sample image
fig, axes = plt.subplots(2, 2)
axes = np.reshape(axes, (4, ))
for i in range(4):
    example = train_loader.__iter__().__next__()[i].numpy().transpose((1, 2, 0))
    axes[i].imshow(example)

```

```

mean = np.array([0.5, 0.5, 0.5])
std = np.array([0.5, 0.5, 0.5])
example = std * example + mean
axes[i].imshow(example)
axes[i].axis('off')
plt.show()

```

```

/home/umhws/anaconda3/envs/eecs442/lib/python3.10/site-packages/torchvision/
transforms/functional.py:153: UserWarning: The given NumPy array is not writ-
able, and PyTorch does not support non-writable tensors. This means writing
to this tensor will result in undefined behavior. You may want to copy the a-
rray to protect its data or make it writable before converting it to a tenso-
r. This type of warning will be suppressed for the rest of this program. (Tr-
iggered internally at ../torch/csrc/utils/tensor_numpy.cpp:206.)
    img = torch.from_numpy(pic.transpose((2, 0, 1))).contiguous()
250
20
torch.Size([5, 3, 256, 256])

```



Step 2: Build Generator and Discriminator

Based on the paper, the architectures of network are as following:

Generator architectures:

U-net encoder:

C64-C128-C256-C512-C512-C512-C512-C512

U-net decoder:

C512-C512-C512-C512-C256-C128-C64-C3

After the last layer in the decoder, a convolution is applied to map to the number of output channels, followed by a Tanh function. As an exception to the above notation, BatchNorm is not applied to the first C64 layer in the encoder. All ReLUs in the encoder are leaky, with slope 0.2, while ReLUs in the decoder are not leaky.

Discriminator architectures

The discriminator architecture is:

C64-C128-C256-C512

After the last layer, a convolution is applied to map to a 1-dimensional output, followed by a Sigmoid function. As an exception to the above notation, BatchNorm is not applied to the first C64 layer. All ReLUs are leaky, with slope 0.2.

We have already implemented the U-net architecture below.

```
In [5]: def normal_init(m, mean, std):
    """
    Helper function. Initialize model parameter with given mean and std.
    """
    if isinstance(m, nn.ConvTranspose2d) or isinstance(m, nn.Conv2d):
        # delete start
        m.weight.data.normal_(mean, std)
        m.bias.data.zero_()
        # delete end
```

```
In [6]: class generator(nn.Module):
    # initializers
    def __init__(self):
        super(generator, self).__init__()

        # Unet generator encoder
        self.conv1 = nn.Conv2d(3, 64, kernel_size=4, stride=2, padding=1)
        self.conv2 = nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1)
        self.bn2 = nn.BatchNorm2d(128)
        self.conv3 = nn.Conv2d(128, 256, kernel_size=4, stride=2, padding=1)
        self.bn3 = nn.BatchNorm2d(256)
        self.conv4 = nn.Conv2d(256, 512, kernel_size=4, stride=2, padding=1)
        self.bn4 = nn.BatchNorm2d(512)
        self.conv5 = nn.Conv2d(512, 512, kernel_size=4, stride=2, padding=1)
        self.bn5 = nn.BatchNorm2d(512)
        self.conv6 = nn.Conv2d(512, 512, kernel_size=4, stride=2, padding=1)
        self.bn6 = nn.BatchNorm2d(512)
        self.conv7 = nn.Conv2d(512, 512, kernel_size=4, stride=2, padding=1)
        self.bn7 = nn.BatchNorm2d(512)
        self.conv8 = nn.Conv2d(512, 512, kernel_size=4, stride=2, padding=1)

        # Unet generator decoder

        self.deconv1 = nn.ConvTranspose2d(512, 512, kernel_size=4, stride=2, padding=1)
        self.bn1 = nn.BatchNorm2d(512)
        self.deconv2 = nn.ConvTranspose2d(1024, 512, kernel_size=4, stride=2, padding=1)
```

```

self.dbn2 = nn.BatchNorm2d(512)
self.deconv3 = nn.ConvTranspose2d(1024, 512, kernel_size=4, stride=2, padding=1)
self.dbn3 = nn.BatchNorm2d(512)
self.deconv4 = nn.ConvTranspose2d(1024, 512, kernel_size=4, stride=2, padding=1)
self.dbn4 = nn.BatchNorm2d(512)
self.deconv5 = nn.ConvTranspose2d(1024, 256, kernel_size=4, stride=2, padding=1)
self.dbn5 = nn.BatchNorm2d(256)
self.deconv6 = nn.ConvTranspose2d(512, 128, kernel_size=4, stride=2, padding=1)
self.dbn6 = nn.BatchNorm2d(128)
self.deconv7 = nn.ConvTranspose2d(256, 64, kernel_size=4, stride=2, padding=1)
self.dbn7 = nn.BatchNorm2d(64)
self.deconv8 = nn.ConvTranspose2d(128, 3, kernel_size=4, stride=2, padding=1)

# weight_init
def weight_init(self, mean, std):
    for m in self._modules:
        normal_init(self._modules[m], mean, std)

# forward method
def forward(self, input):

    # encoding

    e1 = F.leaky_relu(self.conv1(input), negative_slope=0.2)
    e2 = F.leaky_relu(self.bn2(self.conv2(e1)), negative_slope=0.2)
    e3 = F.leaky_relu(self.bn3(self.conv3(e2)), negative_slope=0.2)
    e4 = F.leaky_relu(self.bn4(self.conv4(e3)), negative_slope=0.2)
    e5 = F.leaky_relu(self.bn5(self.conv5(e4)), negative_slope=0.2)
    e6 = F.leaky_relu(self.bn6(self.conv6(e5)), negative_slope=0.2)
    e7 = F.leaky_relu(self.bn7(self.conv7(e6)), negative_slope=0.2)
    e8 = F.leaky_relu(self.conv8(e7), negative_slope=0.2)

    # decoding

    d1 = F.relu(self.dbn1(self.deconv1(e8)))
    d1 = torch.cat([d1, e7], 1)
    d2 = F.relu(self.dbn2(self.deconv2(d1)))
    d2 = torch.cat([d2, e6], 1)
    d3 = F.relu(self.dbn3(self.deconv3(d2)))
    d3 = torch.cat([d3, e5], 1)
    d4 = F.relu(self.dbn4(self.deconv4(d3)))
    d4 = torch.cat([d4, e4], 1)
    d5 = F.relu(self.dbn5(self.deconv5(d4)))
    d5 = torch.cat([d5, e3], 1)
    d6 = F.relu(self.dbn6(self.deconv6(d5)))
    d6 = torch.cat([d6, e2], 1)
    d7 = F.relu(self.dbn7(self.deconv7(d6)))
    d7 = torch.cat([d7, e1], 1)
    d8 = F.tanh(self.deconv8(d7))

    output = d8

    return output

class discriminator(nn.Module):
    # initializers

```

```

def __init__(self):
    super(discriminator, self).__init__()

    self.conv1 = nn.Conv2d(6, 64, kernel_size=4, stride=2, padding=1)
    self.conv2 = nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1)
    self.bn2 = nn.BatchNorm2d(128)
    self.conv3 = nn.Conv2d(128, 256, kernel_size=4, stride=2, padding=1)
    self.bn3 = nn.BatchNorm2d(256)
    self.conv4 = nn.Conv2d(256, 512, kernel_size=4, stride=1, padding=1)
    self.bn4 = nn.BatchNorm2d(512)
    self.conv5 = nn.Conv2d(512, 1, kernel_size=4, stride=1, padding=1)

    # weight_init
def weight_init(self, mean, std):
    for m in self._modules:
        normal_init(self._modules[m], mean, std)

    # forward method
def forward(self, input):

    x1 = F.leaky_relu(self.conv1(input))
    x2 = F.leaky_relu(self.bn2(self.conv2(x1)), negative_slope=0.2)
    x3 = F.leaky_relu(self.bn3(self.conv3(x2)), negative_slope=0.2)
    x4 = F.leaky_relu(self.bn4(self.conv4(x3)), negative_slope=0.2)
    x5 = F.sigmoid(self.conv5(x4))

    x = x5

    return x

```

In [7]: # print out the model summary

```

G = generator().cuda()
D = discriminator().cuda()
summary(G, (3, 256, 256))
summary(D, (6, 256, 256))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 128, 128]	3,136
Conv2d-2	[-1, 128, 64, 64]	131,200
BatchNorm2d-3	[-1, 128, 64, 64]	256
Conv2d-4	[-1, 256, 32, 32]	524,544
BatchNorm2d-5	[-1, 256, 32, 32]	512
Conv2d-6	[-1, 512, 16, 16]	2,097,664
BatchNorm2d-7	[-1, 512, 16, 16]	1,024
Conv2d-8	[-1, 512, 8, 8]	4,194,816
BatchNorm2d-9	[-1, 512, 8, 8]	1,024
Conv2d-10	[-1, 512, 4, 4]	4,194,816
BatchNorm2d-11	[-1, 512, 4, 4]	1,024
Conv2d-12	[-1, 512, 2, 2]	4,194,816
BatchNorm2d-13	[-1, 512, 2, 2]	1,024
Conv2d-14	[-1, 512, 1, 1]	4,194,816
ConvTranspose2d-15	[-1, 512, 2, 2]	4,194,816
BatchNorm2d-16	[-1, 512, 2, 2]	1,024
ConvTranspose2d-17	[-1, 512, 4, 4]	8,389,120
BatchNorm2d-18	[-1, 512, 4, 4]	1,024
ConvTranspose2d-19	[-1, 512, 8, 8]	8,389,120
BatchNorm2d-20	[-1, 512, 8, 8]	1,024
ConvTranspose2d-21	[-1, 512, 16, 16]	8,389,120
BatchNorm2d-22	[-1, 512, 16, 16]	1,024
ConvTranspose2d-23	[-1, 256, 32, 32]	4,194,560
BatchNorm2d-24	[-1, 256, 32, 32]	512
ConvTranspose2d-25	[-1, 128, 64, 64]	1,048,704
BatchNorm2d-26	[-1, 128, 64, 64]	256
ConvTranspose2d-27	[-1, 64, 128, 128]	262,208
BatchNorm2d-28	[-1, 64, 128, 128]	128
ConvTranspose2d-29	[-1, 3, 256, 256]	6,147

Total params: 54,419,459

Trainable params: 54,419,459

Non-trainable params: 0

Input size (MB): 0.75

Forward/backward pass size (MB): 54.82

Params size (MB): 207.59

Estimated Total Size (MB): 263.16

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 128, 128]	6,208
Conv2d-2	[-1, 128, 64, 64]	131,200
BatchNorm2d-3	[-1, 128, 64, 64]	256
Conv2d-4	[-1, 256, 32, 32]	524,544
BatchNorm2d-5	[-1, 256, 32, 32]	512
Conv2d-6	[-1, 512, 31, 31]	2,097,664
BatchNorm2d-7	[-1, 512, 31, 31]	1,024
Conv2d-8	[-1, 1, 30, 30]	8,193

Total params: 2,769,601

Trainable params: 2,769,601

```
Non-trainable params: 0
```

```
-----  
Input size (MB): 1.50  
Forward/backward pass size (MB): 27.51  
Params size (MB): 10.57  
Estimated Total Size (MB): 39.58  
-----
```

```
In [8]: D
```

```
Out[8]: discriminator(  
    (conv1): Conv2d(6, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    (conv2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
    1))  
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv3): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
    1))  
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv4): Conv2d(256, 512, kernel_size=(4, 4), stride=(1, 1), padding=(1,  
    1))  
    (bn4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv5): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), padding=(1,  
    1))  
)
```

```
In [9]: G
```

```
Out[9]: generator(  
    (conv1): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
    (conv2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
1))  
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv3): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
1))  
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv4): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
1))  
    (bn4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv5): Conv2d(512, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
1))  
    (bn5): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv6): Conv2d(512, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
1))  
    (bn6): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv7): Conv2d(512, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
1))  
    (bn7): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runni  
ng_stats=True)  
    (conv8): Conv2d(512, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1,  
1))  
    (deconv1): ConvTranspose2d(512, 512, kernel_size=(4, 4), stride=(2, 2), p  
adding=(1, 1))  
    (dbn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runn  
ing_stats=True)  
    (deconv2): ConvTranspose2d(1024, 512, kernel_size=(4, 4), stride=(2, 2),  
padding=(1, 1))  
    (dbn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runn  
ing_stats=True)  
    (deconv3): ConvTranspose2d(1024, 512, kernel_size=(4, 4), stride=(2, 2),  
padding=(1, 1))  
    (dbn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runn  
ing_stats=True)  
    (deconv4): ConvTranspose2d(1024, 512, kernel_size=(4, 4), stride=(2, 2),  
padding=(1, 1))  
    (dbn4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_runn  
ing_stats=True)  
    (deconv5): ConvTranspose2d(1024, 256, kernel_size=(4, 4), stride=(2, 2),  
padding=(1, 1))  
    (dbn5): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_runn  
ing_stats=True)  
    (deconv6): ConvTranspose2d(512, 128, kernel_size=(4, 4), stride=(2, 2), p  
adding=(1, 1))  
    (dbn6): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_runn  
ing_stats=True)  
    (deconv7): ConvTranspose2d(256, 64, kernel_size=(4, 4), stride=(2, 2), pa  
dding=(1, 1))  
    (dbn7): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_runn  
ing_stats=True)
```

```
(deconv8): ConvTranspose2d(128, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
```

Step 3: Train

In this section, We will complete the training function. Then train our model. Below are some helper function that can be used. We will guide you through the implementation.

```
In [10]: # Helper function for showing result.
def process_image(img):
    return (img.cpu().data.numpy().transpose(1, 2, 0) + 1) / 2

def show_result(G, x_, y_, num_epoch):
    predict_images = G(x_)

    fig, ax = plt.subplots(x_.size()[0], 3, figsize=(6,10))

    for i in range(x_.size()[0]):
        ax[i, 0].get_xaxis().set_visible(False)
        ax[i, 0].get_yaxis().set_visible(False)
        ax[i, 1].get_xaxis().set_visible(False)
        ax[i, 1].get_yaxis().set_visible(False)
        ax[i, 2].get_xaxis().set_visible(False)
        ax[i, 2].get_yaxis().set_visible(False)
        ax[i, 0].cla()
        ax[i, 0].imshow(process_image(x_[i]))
        ax[i, 1].cla()
        ax[i, 1].imshow(process_image(predict_images[i]))
        ax[i, 2].cla()
        ax[i, 2].imshow(process_image(y_[i]))

    plt.tight_layout()
    label_epoch = 'Epoch {0}'.format(num_epoch)
    fig.text(0.5, 0, label_epoch, ha='center')
    label_input = 'Input'
    fig.text(0.18, 1, label_input, ha='center')
    label_output = 'Output'
    fig.text(0.5, 1, label_output, ha='center')
    label_truth = 'Ground truth'
    fig.text(0.81, 1, label_truth, ha='center')

    plt.show()

# Helper function for counting number of trainable parameters.
def count_params(model):
    ...
    Counts the number of trainable parameters in PyTorch.
    Args:
        model: PyTorch model.
    Returns:
        num_params: int, number of trainable parameters.
    ...
```

```
num_params = sum([item.numel() for item in model.parameters() if item.requires_grad])
return num_params
```

Objective Functions: Conditional Pix2Pix GAN

Given a generator G and a discriminator D , the loss function / objective functions to be minimized are given by

$$\mathcal{L}_{cGAN}(G, D) = \frac{1}{N} \left(\sum_{i=1}^N \log D(x_i, y_i) \right)$$

- $\sum_{i=1}^N \log (1 - D(G(x_i), y_i))$

where (x_i, y_i) refers to the pair to the ground-truth input-output pair and $G(x_i)$ refers to the image translated by the Generator.

$$\mathcal{L}_{L1}(G, D) = \frac{1}{N} \sum_{i=1}^N \|y - G(x_i)\|_1$$

The final objective is just a combination of these objectives.

$$\mathcal{L}_{final}(G, D) = \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G, D)$$

$$G^* = \underset{G}{\operatorname{argmin}} \underset{D}{\operatorname{max}} \mathcal{L}_{final}(G, D)$$

You would be implementing these objectives using the `nn.BCELoss` and `nn.L1Loss` as given below.

```
In [11]: # Hint: you could use following loss to complete following function
BCE_loss = nn.BCELoss().cuda()
L1_loss = nn.L1Loss().cuda()

def train(G, D, num_epochs = 20):
    hist_D_losses = []
    hist_G_losses = []
    hist_G_L1_losses = []
    #####
    # TODO: Add Adam optimizer to generator and discriminator
    # You will use lr=0.0002, beta=0.5, beta2=0.999
    #####
    G_optimizer = optim.Adam(G.parameters(), lr=2e-4, betas=(0.5, 0.999))
    D_optimizer = optim.Adam(D.parameters(), lr=2e-4, betas=(0.5, 0.999))

    #####
    # END OF YOUR CODE
    #####
    print('training start!')
    start_time = time.time()
    for epoch in range(num_epochs):
        print('Start training epoch %d' % (epoch + 1))
        D_losses = []
```

```

G_losses = []
epoch_start_time = time.time()
num_iter = 0
for x_ in train_loader:
    y_ = x_[:, :, :, img_size:]
    x_ = x_[:, :, :, 0:img_size]

    x_, y_ = x_.cuda(), y_.cuda()
    ##### Implement training code for the discriminator.
    # Recall that the loss is the mean of the loss for real images and fake
    # images, and made by some calculations with zeros and ones
    # We have defined the BCE_loss, which you might would like to use.
    #
    # NOTE: While training the Discriminator, the output of the generator
    # must be detached from the computational graph. Refer to the method
    # torch.Tensor.detach()
    #####
    N = x_.shape[0]
    # Generate data
    fake_data = G.forward(x_).detach()

    #1. Train the discriminator
    # D real data BCE loss
    D_real_preds = D.forward(torch.cat((x_, y_), dim=1))
    D_y_real = torch.ones_like(D_real_preds)
    # D_real_loss = torch.sum(torch.log(D_real_preds))
    D_real_loss = BCE_loss(D_real_preds, D_y_real)

    # D fake data BCE loss
    D_fake_preds = D.forward(torch.cat((x_, fake_data), dim=1))
    D_y_fake = torch.zeros_like(D_fake_preds)
    # D_fake_loss = torch.sum(torch.log(1 - D_fake_preds))
    D_fake_loss = BCE_loss(D_fake_preds, D_y_fake)

    # D loss
    loss_D = D_real_loss + D_fake_loss

    # Train D
    D_optimizer.zero_grad()
    loss_D.backward()
    D_optimizer.step()

    #####
    # END OF YOUR CODE
    #####
    #####
    # TODO: Implement training code for the Generator.
    #####
    # 1. Train the generator
    # 2. Append the losses to the lists 'hist_G_L1_losses' and 'hist_G_losses'
    # (Only append the data to the list, not the complete tensor, refer
    # to torch.Tensor.item()).

```

```

# Generate data
fake_data = G.forward(x_)

# 1. Train the generator
# G BCE loss
G_fake_preds = D.forward(torch.cat((x_, fake_data), dim=1))
G_y_fake = torch.zeros_like(G_fake_preds)
G_bce_loss = BCE_loss(G_fake_preds, G_y_fake)

# G l1 loss
G_l1_loss = L1_loss(fake_data, y_)

# G loss
lamb = 100
loss_G = G_bce_loss + lamb * G_l1_loss

# Train G
G_optimizer.zero_grad()
loss_G.backward()
G_optimizer.step()

# 2. Append the losses to the lists 'hist_G_L1_losses' and 'hist_D_losses'
# (Only append the data to the list, not the complete tensor, refer
# torch.Tensor.item()).
hist_G_losses.append(G_bce_loss.detach().item())
hist_G_L1_losses.append(G_l1_loss.detach().item())
#####
# END OF YOUR CODE
#####

D_losses.append(loss_D.detach().item())
hist_D_losses.append(loss_D.detach().item())
G_losses.append(loss_G)
num_iter += 1

epoch_end_time = time.time()
per_epoch_ptime = epoch_end_time - epoch_start_time

print('[%d/%d] - using time: %.2f seconds' % ((epoch + 1), num_epochs, per_epoch_ptime))
print('loss of discriminator D: %.3f' % (torch.mean(torch.FloatTensor(D_losses))))
print('loss of generator G: %.3f' % (torch.mean(torch.FloatTensor(G_losses))))
if epoch == 0 or (epoch + 1) % 5 == 0:
    with torch.no_grad():
        show_result(G, fixed_x_, fixed_y_, (epoch+1))

end_time = time.time()
total_ptime = end_time - start_time

return hist_D_losses, hist_G_losses, hist_G_L1_losses

```

In this part, train your model with $\lambda = 100$ with at least 20 epochs.

In [12]: # Define network

```
G_100 = generator()
D_100 = discriminator()
G_100.weight_init(mean=0.0, std=0.02)
D_100.weight_init(mean=0.0, std=0.02)
G_100.cuda()
D_100.cuda()
G_100.train()
D_100.train()
```

Out[12]: discriminator

```
(conv1): Conv2d(6, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
(conv2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
(bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv3): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1,
1))
(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv4): Conv2d(256, 512, kernel_size=(4, 4), stride=(1, 1), padding=(1,
1))
(bn4): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv5): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), padding=(1,
1))
)
```

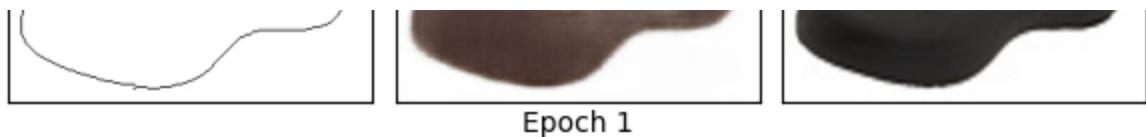
In [13]: #training, you will be expecting 1-2 minutes per epoch.

```
# TODO: change num_epochs if you want
```

```
hist_D_100_losses, hist_G_100_BCE_losses, hist_G_100_L1_losses = train(G=G_1)

training start!
Start training epoch 1
[1/20] - using time: 33.23 seconds
loss of discriminator D: 0.521
loss of generator G: 25.637
```





Start training epoch 2
[2/20] - using time: 30.38 seconds
loss of discriminator D: 0.233
loss of generator G: 18.436
Start training epoch 3
[3/20] - using time: 30.51 seconds
loss of discriminator D: 0.117
loss of generator G: 17.228
Start training epoch 4
[4/20] - using time: 30.43 seconds
loss of discriminator D: 0.188
loss of generator G: 16.127
Start training epoch 5
[5/20] - using time: 30.52 seconds
loss of discriminator D: 0.041
loss of generator G: 15.267





```
Start training epoch 6
[6/20] - using time: 30.44 seconds
loss of discriminator D: 0.046
loss of generator G: 13.856
Start training epoch 7
[7/20] - using time: 30.49 seconds
loss of discriminator D: 0.005
loss of generator G: 13.028
Start training epoch 8
[8/20] - using time: 30.37 seconds
loss of discriminator D: 0.097
loss of generator G: 12.411
Start training epoch 9
[9/20] - using time: 30.57 seconds
loss of discriminator D: 0.086
loss of generator G: 11.618
Start training epoch 10
[10/20] - using time: 30.50 seconds
loss of discriminator D: 0.006
loss of generator G: 11.079
```





```
Start training epoch 11
[11/20] - using time: 30.36 seconds
loss of discriminator D: 0.033
loss of generator G: 10.787
Start training epoch 12
[12/20] - using time: 30.41 seconds
loss of discriminator D: 0.197
loss of generator G: 10.169
Start training epoch 13
[13/20] - using time: 30.40 seconds
loss of discriminator D: 0.060
loss of generator G: 9.624
Start training epoch 14
[14/20] - using time: 30.40 seconds
loss of discriminator D: 0.066
loss of generator G: 9.171
Start training epoch 15
[15/20] - using time: 30.53 seconds
loss of discriminator D: 0.004
loss of generator G: 8.806
```





Epoch 15

```
Start training epoch 16
[16/20] - using time: 30.52 seconds
loss of discriminator D: 0.003
loss of generator G: 8.398
Start training epoch 17
[17/20] - using time: 30.42 seconds
loss of discriminator D: 0.005
loss of generator G: 8.094
Start training epoch 18
[18/20] - using time: 30.77 seconds
loss of discriminator D: 0.004
loss of generator G: 7.857
Start training epoch 19
[19/20] - using time: 30.47 seconds
loss of discriminator D: 0.003
loss of generator G: 7.606
Start training epoch 20
[20/20] - using time: 30.52 seconds
loss of discriminator D: 0.216
loss of generator G: 7.544
```





```
In [14]: !mkdir models
torch.save(G_100.state_dict(), './models/G_100.pth')
torch.save(D_100.state_dict(), './models/D_100.pth')
```

mkdir: cannot create directory 'models': File exists

The following cell saves the trained model parameters to your Google Drive so you could reuse those parameters later without retraining.

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
!cp "./models/" -r "/content/drive/My Drive/"
```

Mounted at /content/drive

Step 4: Visualization

Please plot the generator BCE and L1 losses, as well as the discriminator loss. For this, please use $\lambda = 100$, and use 3 separate plots. We have provided the code for you and you only need to run the code below.

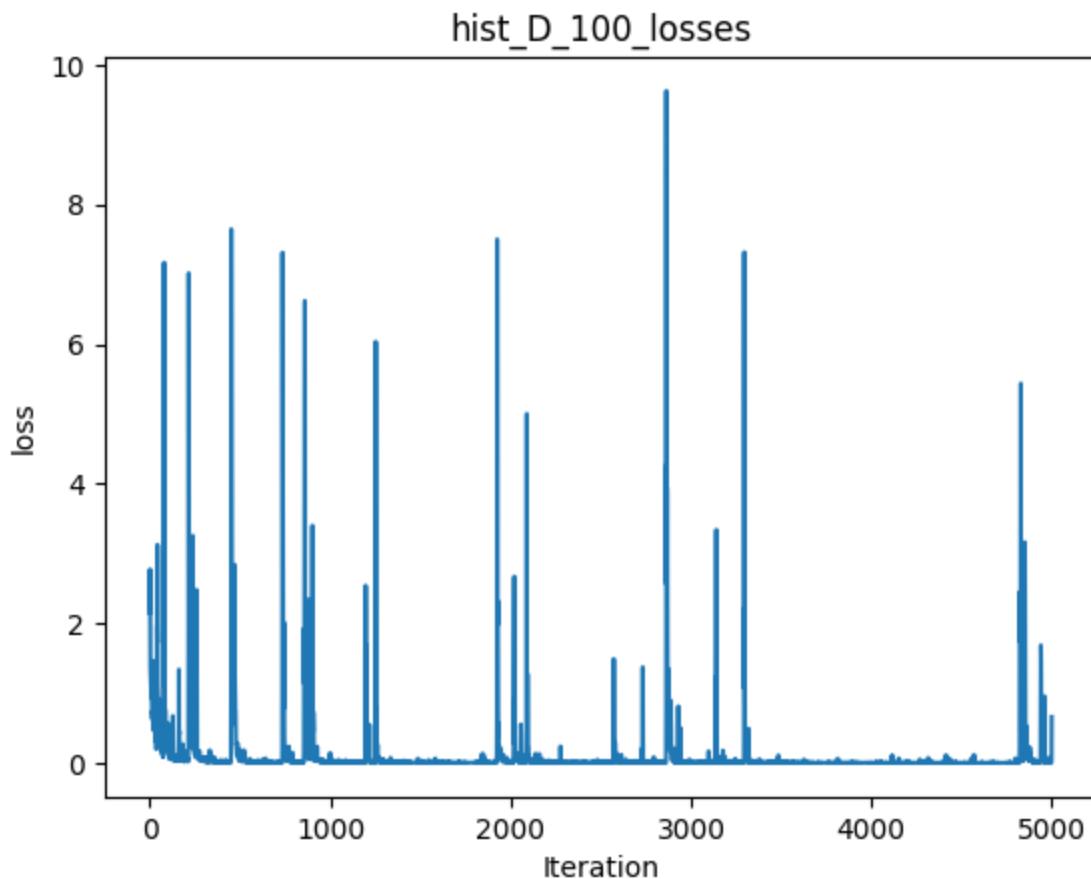
```
In [15]: # Plot the G/D loss history (y axis) vs. Iteration (x axis)      #
# You will have three plots, with hist_D_100_losses,                      #
# hist_G_100_BCE_losses, hist_G_100_L1_losses respectively.                 #

# hist_D_100_losses
plt.figure()
plt.plot(range(len(hist_D_100_losses)), torch.tensor(hist_D_100_losses, device='cpu'))
# plt.legend('hist_D_100_losses')
plt.xlabel('Iteration')
plt.ylabel('loss')
plt.title('hist_D_100_losses')
plt.show()

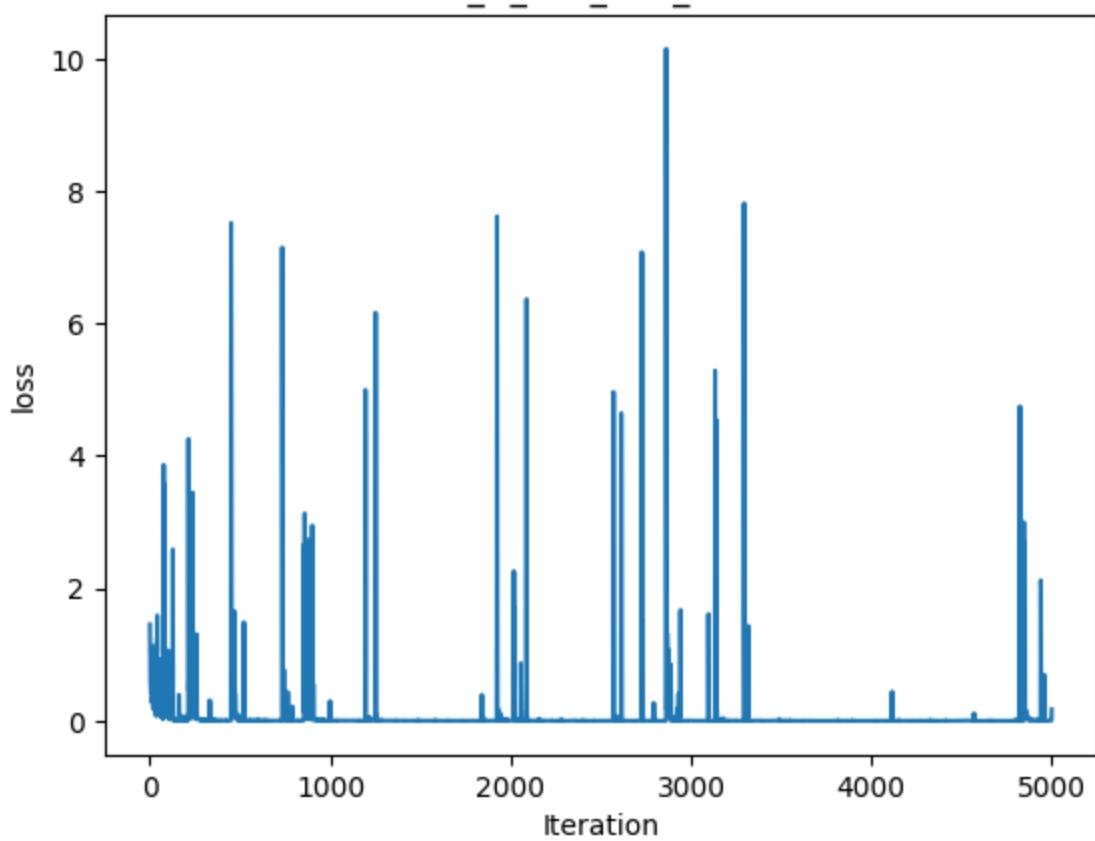
# hist_G_100_BCE_losses
plt.figure()
plt.plot(range(len(hist_G_100_BCE_losses)), torch.tensor(hist_G_100_BCE_losses, device='cpu'))
# plt.legend('hist_G_100_BCE_losses')
plt.xlabel('Iteration')
plt.ylabel('loss')
plt.title('hist_G_100_BCE_losses')
plt.show()

# hist_G_100_L1_losses
plt.figure()
plt.plot(range(len(hist_G_100_L1_losses)), torch.tensor(hist_G_100_L1_losses, device='cpu'))
# plt.legend('hist_G_100_L1_losses')
```

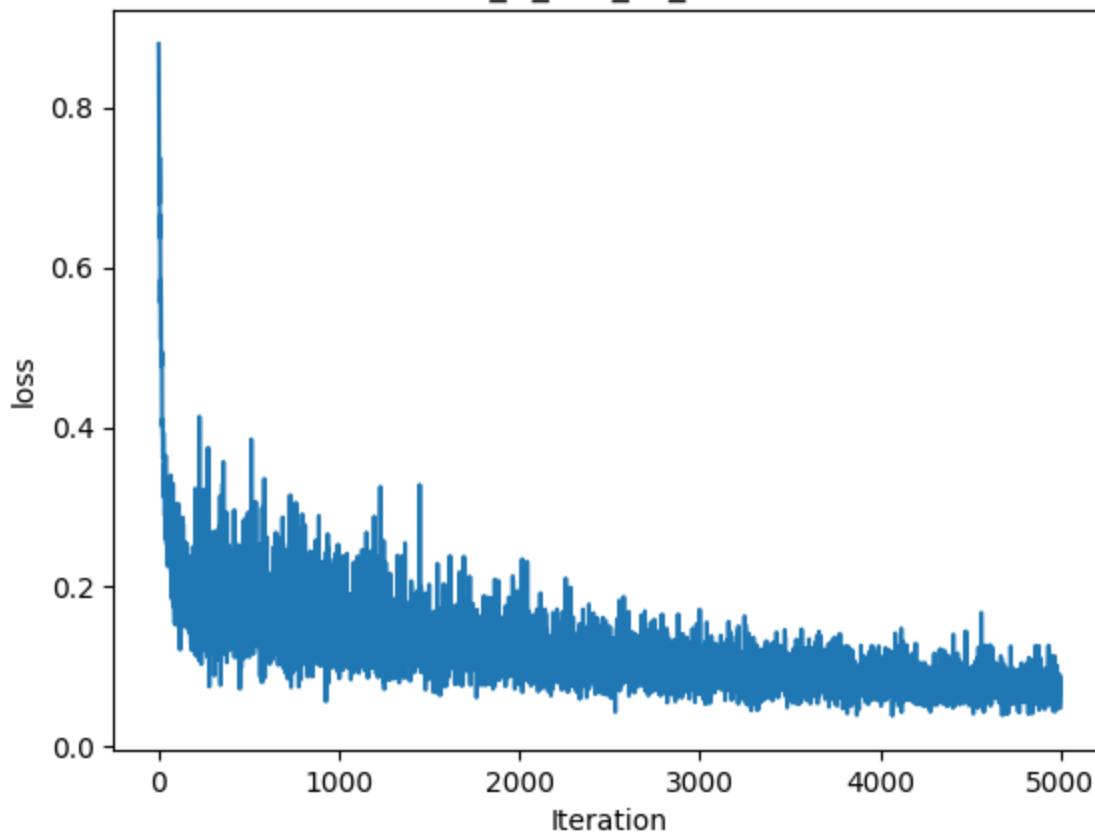
```
plt.xlabel('Iteration')
plt.ylabel('loss')
plt.title('hist_G_100_L1_losses')
plt.show()
```



hist_G_100_BCE_losses



hist_G_100_L1_losses



Convert Notebook to PDF

Alternative if the cell below doesn't work.

```
In [ ]: import os
from google.colab import drive
from google.colab import files

drive.mount_point = '/content/drive/'
drive.mount(drive.mount_point)
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force_remount=True).

```
In [ ]: # generate pdf
# Please provide the full path of the notebook file below
# Important: make sure that your file name does not contain spaces!

# Ex: notebookpath = '/content/drive/My Drive/Colab Notebooks/EECS_442_PS4_F'
notebookpath = '/content/drive/My Drive/Colab Notebooks/EECS_442-504_PS6_FA_'

file_name = notebookpath.split('/')[-1]
get_ipython().system("apt update && apt install texlive-xetex texlive-fonts-recommended")
get_ipython().system("jupyter nbconvert --to PDF {}".format(notebookpath.replace('.ipynb', '.pdf')))

files.download(notebookpath.split('.')[0]+'.pdf')
```

Homework 5: Diffusion Models

Run the following code to setup the necessary requirements

```
In [ ]: from google.colab import drive  
drive.mount('/content/drive')
```

```
In [ ]: # TODO: Fill in the Google Drive path where you uploaded the assignment  
# Example: If you create a EECS442 folder and put all the files under HW5 fo  
GOOGLE_DRIVE_PATH_AFTER_MYDRIVE = 'EECS442/HW5'
```

```
In [1]: %load_ext autoreload  
%autoreload 2
```

```
In [ ]: import os  
import sys  
  
GOOGLE_DRIVE_PATH = os.path.join('drive', 'MyDrive', GOOGLE_DRIVE_PATH_AFT  
sys.path.append(GOOGLE_DRIVE_PATH)
```

```
In [ ]: print(GOOGLE_DRIVE_PATH)
```

You need to change your working directory.

```
In [ ]: %cd /content/drive/MyDrive/EECS442/HW5
```

```
In [ ]: !pip install certifi>=2022.9.14  
!pip install charset-normalizer>=2.1.1  
!pip install contourpy>=1.0.5  
!pip install cycler>=0.11.0  
!pip install fonttools>=4.37.2  
!pip install idna>=3.4  
!pip install kiwisolver>=1.4.4  
!pip install matplotlib>=3.6.0  
!pip install numpy>=1.23.3  
!pip install packaging>=21.3  
!pip install Pillow>=9.2.0  
!pip install pyparsing>=3.0.9  
!pip install python-dateutil>=2.8.2  
!pip install PyYAML>=6.0  
!pip install requests>=2.28.1  
!pip install scipy>=1.9.1  
!pip install six>=1.16.0  
!pip install tqdm>=4.64.1  
!pip install typing-extensions>=4.3.0  
!pip install urllib3>=1.26.12
```

```
In [ ]: !nvidia-smi
```

Task 1: Unconditional Sampling with DDPM

Setup

```
In [2]: from functools import partial
import os
import argparse
import yaml

import torch
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

from util.logger import get_logger
# GPUs are preferred
logger = get_logger()
device_str = f"cuda:0" if torch.cuda.is_available() else 'cpu'
logger.info(f"Device set to {device_str}.")
device = torch.device(device_str)

# set output directory
save_dir = 'results'
ddpm_out_path = os.path.join(save_dir, 'uncond_ddpm')
os.makedirs(ddpm_out_path, exist_ok=True)
for img_dir in ['input', 'output', 'progress']:
    os.makedirs(os.path.join(ddpm_out_path, img_dir), exist_ok=True)
```

2024-04-09 14:31:59,649 [DPS] >> Device set to cuda:0.

In this task, you will implement the sampling Algorithm proposed in Denoising Diffusion Probabilistic Models(DDPM) paper as shown below:

(1) Now let's implement the variance schedule. As you can see in the DDPM sampling algorithm, We will need α_t for each timestep t . α_t is a notaion for $1 - \beta_t$, where β_t is the true variance that increases from $t = 1$ to $t = T$. There are many different variance schedules such as linear schedule and cosine schedule. Follow the instruction in `guided_diffusion/simple_diffusion.py` to implement `get_named_beta_schedule()`.

Cosine schedule is proposed by [iDDPM](#). You can find the detailed motivation in the paper. The calculation of β depends on α , the cumulated product of α is defined as

$$\bar{\alpha}_t = \frac{f(t)}{f(0)}$$

, where

$$f(t) = \cos\left(\frac{t/T+s}{1+s} \cdot \frac{\pi}{2}\right)^2$$

We use small $s = 0.008$ such that $\sqrt{\beta_0}$ was slightly smaller than the pixel bin size 1/127.5. According to the definition of α_t , we can then get β_t as

$$\beta_t = 1 - \frac{\bar{\alpha}_t}{\bar{\alpha}_{t-1}}$$

Also, clip β_t to be no larger than 0.999 to prevent singularities at the end of the diffusion process.

Run the following code to see your output.

```
In [3]: from guided_diffusion.simple_diffusion import get_named_beta_schedule
import numpy as np

num_steps = 1000
schedule_name = 'cosine'

print('Cosine Error: ', np.sum(get_named_beta_schedule(schedule_name, num_st

```

/home/umhws/anaconda3/envs/eecs442/lib/python3.10/site-packages/tqdm/auto.p
y:21: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
Cosine Error: -6.661338147750939e-16

(2) Now you have implemented your variance schedule $\{\beta_1, \beta_2, \dots, \beta_T\}$. In practice, we use α_t and accumulated product $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$. **Follow the instruction to complete `__init__()` of DDPMDiffusion** to compute the needed values that will be used during the sampling process. They only need to be calculated once during initialization and we can directly access them later during sampling.

Hint: You can use `np.cumprod()` to calculate the cumulated product.

Let's now create the model.

Before you run the code below, make sure that you have downloaded the pretrained model `ffhq_10M.pt` and put it under `model` directory.

```
In [4]: from guided_diffusion.unet import create_model
from data.dataloader import get_dataset, get_dataloader

# Here is the model configuration of the Pretrained UNet model that we will
# This configuration should be consistent with the pretrained model, so you
# You can find the detailed definition of the UNet in guided_diffusion/unet.

model_config = {
    'image_size': 256,
    'num_channels': 128,
```

```
'num_res_blocks': 1,
'channel_mult': '',
'learn_sigma': True,
'class_cond': False,
'use_checkpoint': False,
'attention_resolutions': 16,
'num_heads': 4,
'num_head_channels': 64,
'num_heads_upsample': -1,
'use_scale_shift_norm': True,
'dropout': 0.0,
'resblock_updown': True,
'use_fp16': False,
'use_new_attention_order': False,
'model_path': 'models/ffhq_10m.pt'
}

# Load model
ddpm_beta = get_named_beta_schedule('linear', 1000)

model = create_model(betas=ddpm_beta, **model_config)
model = model.to(device)
model.eval() # Set the model to the evaluation mode as we don't need to tra
```

pretrained model loaded!

```
Out[4]: UNetModel(  
    (time_embed): Sequential(  
        (0): Linear(in_features=128, out_features=512, bias=True)  
        (1): SiLU()  
        (2): Linear(in_features=512, out_features=512, bias=True)  
    )  
    (input_blocks): ModuleList(  
        (0): TimestepEmbedSequential(  
            (0): Conv2d(3, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,  
                1))  
        )  
        (1): TimestepEmbedSequential(  
            (0): ResBlock(  
                (in_layers): Sequential(  
                    (0): GroupNorm32(32, 128, eps=1e-05, affine=True)  
                    (1): SiLU()  
                    (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=  
                        (1, 1))  
                )  
                (h_upd): Identity()  
                (x_upd): Identity()  
                (emb_layers): Sequential(  
                    (0): SiLU()  
                    (1): Linear(in_features=512, out_features=256, bias=True)  
                )  
                (out_layers): Sequential(  
                    (0): GroupNorm32(32, 128, eps=1e-05, affine=True)  
                    (1): SiLU()  
                    (2): Dropout(p=0.0, inplace=False)  
                    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=  
                        (1, 1))  
                )  
                (skip_connection): Identity()  
            )  
        )  
        (2): TimestepEmbedSequential(  
            (0): ResBlock(  
                (in_layers): Sequential(  
                    (0): GroupNorm32(32, 128, eps=1e-05, affine=True)  
                    (1): SiLU()  
                    (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=  
                        (1, 1))  
                )  
                (h_upd): Downsample(  
                    (op): AvgPool2d(kernel_size=2, stride=2, padding=0)  
                )  
                (x_upd): Downsample(  
                    (op): AvgPool2d(kernel_size=2, stride=2, padding=0)  
                )  
                (emb_layers): Sequential(  
                    (0): SiLU()  
                    (1): Linear(in_features=512, out_features=256, bias=True)  
                )  
                (out_layers): Sequential(  
                    (0): GroupNorm32(32, 128, eps=1e-05, affine=True)  
                    (1): SiLU()  
                )  
            )  
        )  
    )  
)
```

```

        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (skip_connection): Identity()
)
)
(3): TimestepEmbedSequential(
(0): ResBlock(
    (in_layers): Sequential(
        (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (h_upd): Identity()
    (x_upd): Identity()
    (emb_layers): Sequential(
        (0): SiLU()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (out_layers): Sequential(
        (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (skip_connection): Identity()
)
)
(4): TimestepEmbedSequential(
(0): ResBlock(
    (in_layers): Sequential(
        (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (h_upd): Downsample(
        (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
    )
    (x_upd): Downsample(
        (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
    )
    (emb_layers): Sequential(
        (0): SiLU()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (out_layers): Sequential(
        (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
)
)

```

```

        (skip_connection): Identity()
    )
)
(5): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=512, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Conv2d(128, 256, kernel_size=(1, 1), stride=(1,
1))
    )
)
(6): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Downsample(
            (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
        )
        (x_upd): Downsample(
            (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
        )
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=512, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Identity()
    )
)

```

```

(7): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=512, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Identity()
    )
)
(8): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Downsample(
            (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
        )
        (x_upd): Downsample(
            (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
        )
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=512, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Identity()
    )
)
(9): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)

```

```

        (1): SiLU()
        (2): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (h_upd): Identity()
    (x_upd): Identity()
    (emb_layers): Sequential(
        (0): SiLU()
        (1): Linear(in_features=512, out_features=1024, bias=True)
    )
    (out_layers): Sequential(
        (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (skip_connection): Conv2d(256, 512, kernel_size=(1, 1), stride=(1,
1))
)
(1): AttentionBlock(
    (norm): GroupNorm32(32, 512, eps=1e-05, affine=True)
    (qkv): Conv1d(512, 1536, kernel_size=(1,), stride=(1,))
    (attention): QKVAttentionLegacy()
    (proj_out): Conv1d(512, 512, kernel_size=(1,), stride=(1,))
)
)
(10): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Downsample(
            (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
        )
        (x_upd): Downsample(
            (op): AvgPool2d(kernel_size=2, stride=2, padding=0)
        )
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=1024, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Identity()
    )
)
(11): TimestepEmbedSequential(

```

```
(0): ResBlock(
    (in_layers): Sequential(
        (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (h_upd): Identity()
    (x_upd): Identity()
    (emb_layers): Sequential(
        (0): SiLU()
        (1): Linear(in_features=512, out_features=1024, bias=True)
    )
    (out_layers): Sequential(
        (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (skip_connection): Identity()
)
)
)
(middle_block): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=1024, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Identity()
    )
)
)
(1): AttentionBlock(
    (norm): GroupNorm32(32, 512, eps=1e-05, affine=True)
    (qkv): Conv1d(512, 1536, kernel_size=(1,), stride=(1,))
    (attention): QKVAAttentionLegacy()
    (proj_out): Conv1d(512, 512, kernel_size=(1,), stride=(1,))
)
)
(2): ResBlock(
    (in_layers): Sequential(
        (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
    )
)
```

```

        (1): SiLU()
        (2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (h_upd): Identity()
    (x_upd): Identity()
    (emb_layers): Sequential(
        (0): SiLU()
        (1): Linear(in_features=512, out_features=1024, bias=True)
    )
    (out_layers): Sequential(
        (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (skip_connection): Identity()
)
)
(output_blocks): ModuleList(
    (0): TimestepEmbedSequential(
        (0): ResBlock(
            (in_layers): Sequential(
                (0): GroupNorm32(32, 1024, eps=1e-05, affine=True)
                (1): SiLU()
                (2): Conv2d(1024, 512, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1))
            )
            (h_upd): Identity()
            (x_upd): Identity()
            (emb_layers): Sequential(
                (0): SiLU()
                (1): Linear(in_features=512, out_features=1024, bias=True)
            )
            (out_layers): Sequential(
                (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
                (1): SiLU()
                (2): Dropout(p=0.0, inplace=False)
                (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
            )
            (skip_connection): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1,
1))
        )
    )
)
(1): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 1024, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(1024, 512, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
    )
)

```

```

        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=1024, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1,
1))
    )
    (1): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Upsample()
        (x_upd): Upsample()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=1024, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Identity()
    )
)
(2): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 1024, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(1024, 512, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=1024, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=

```

```

(1, 1))
)
(skip_connection): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1,
1))
)
(1): AttentionBlock(
    (norm): GroupNorm32(32, 512, eps=1e-05, affine=True)
    (qkv): Conv1d(512, 1536, kernel_size=(1,), stride=(1,))
    (attention): QKVAttentionLegacy()
    (proj_out): Conv1d(512, 512, kernel_size=(1,), stride=(1,))
)
)
(3): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 768, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(768, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=1024, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Conv2d(768, 512, kernel_size=(1, 1), stride=(1,
1))
    )
    (1): AttentionBlock(
        (norm): GroupNorm32(32, 512, eps=1e-05, affine=True)
        (qkv): Conv1d(512, 1536, kernel_size=(1,), stride=(1,))
        (attention): QKVAttentionLegacy()
        (proj_out): Conv1d(512, 512, kernel_size=(1,), stride=(1,))
    )
    (2): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Upsample()
        (x_upd): Upsample()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=1024, bias=True)
        )
        (out_layers): Sequential(

```

```

(0): GroupNorm32(32, 512, eps=1e-05, affine=True)
(1): SiLU()
(2): Dropout(p=0.0, inplace=False)
(3): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
(skip_connection): Identity()
)
)
(4): TimestepEmbedSequential(
(0): ResBlock(
(in_layers): Sequential(
(0): GroupNorm32(32, 768, eps=1e-05, affine=True)
(1): SiLU()
(2): Conv2d(768, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
(h_upd): Identity()
(x_upd): Identity()
(emb_layers): Sequential(
(0): SiLU()
(1): Linear(in_features=512, out_features=512, bias=True)
)
(out_layers): Sequential(
(0): GroupNorm32(32, 256, eps=1e-05, affine=True)
(1): SiLU()
(2): Dropout(p=0.0, inplace=False)
(3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
(skip_connection): Conv2d(768, 256, kernel_size=(1, 1), stride=(1,
1))
)
)
)
(5): TimestepEmbedSequential(
(0): ResBlock(
(in_layers): Sequential(
(0): GroupNorm32(32, 512, eps=1e-05, affine=True)
(1): SiLU()
(2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
(h_upd): Identity()
(x_upd): Identity()
(emb_layers): Sequential(
(0): SiLU()
(1): Linear(in_features=512, out_features=512, bias=True)
)
(out_layers): Sequential(
(0): GroupNorm32(32, 256, eps=1e-05, affine=True)
(1): SiLU()
(2): Dropout(p=0.0, inplace=False)
(3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
(skip_connection): Conv2d(512, 256, kernel_size=(1, 1), stride=(1,
1))
)
)
)

```

```

    1))
    )
    (1): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
                (1, 1))
        )
        (h_upd): Upsample()
        (x_upd): Upsample()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=512, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
                (1, 1))
        )
        (skip_connection): Identity()
    )
)
(6): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 512, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=
                (1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=512, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
                (1, 1))
        )
        (skip_connection): Conv2d(512, 256, kernel_size=(1, 1), stride=(1,
                1))
    )
)
(7): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 384, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=

```

```

(1, 1))
)
(h_upd): Identity()
(x_upd): Identity()
(emb_layers): Sequential(
    (0): SiLU()
    (1): Linear(in_features=512, out_features=512, bias=True)
)
(out_layers): Sequential(
    (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
    (1): SiLU()
    (2): Dropout(p=0.0, inplace=False)
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
(skip_connection): Conv2d(384, 256, kernel_size=(1, 1), stride=(1,
1))
)
(1): ResBlock(
    (in_layers): Sequential(
        (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
)
(h_upd): Upsample()
(x_upd): Upsample()
(emb_layers): Sequential(
    (0): SiLU()
    (1): Linear(in_features=512, out_features=512, bias=True)
)
(out_layers): Sequential(
    (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
    (1): SiLU()
    (2): Dropout(p=0.0, inplace=False)
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
)
(skip_connection): Identity()
)
)
)
(8): TimestepEmbedSequential(
(0): ResBlock(
    (in_layers): Sequential(
        (0): GroupNorm32(32, 384, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Conv2d(384, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
)
)
(h_upd): Identity()
(x_upd): Identity()
(emb_layers): Sequential(
    (0): SiLU()
    (1): Linear(in_features=512, out_features=256, bias=True)
)
(out_layers): Sequential(

```

```

        (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (skip_connection): Conv2d(384, 128, kernel_size=(1, 1), stride=(1,
1))
)
)
)
(9): TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=256, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Conv2d(256, 128, kernel_size=(1, 1), stride=(1,
1))
)
)
(1): ResBlock(
    (in_layers): Sequential(
        (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (h_upd): Upsample()
    (x_upd): Upsample()
    (emb_layers): Sequential(
        (0): SiLU()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (out_layers): Sequential(
        (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
        (1): SiLU()
        (2): Dropout(p=0.0, inplace=False)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    )
    (skip_connection): Identity()
)
)

```

```
(10-11): 2 x TimestepEmbedSequential(
    (0): ResBlock(
        (in_layers): Sequential(
            (0): GroupNorm32(32, 256, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (h_upd): Identity()
        (x_upd): Identity()
        (emb_layers): Sequential(
            (0): SiLU()
            (1): Linear(in_features=512, out_features=256, bias=True)
        )
        (out_layers): Sequential(
            (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
            (1): SiLU()
            (2): Dropout(p=0.0, inplace=False)
            (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
        )
        (skip_connection): Conv2d(256, 128, kernel_size=(1, 1), stride=(1,
1))
    )
)
)
(out): Sequential(
    (0): GroupNorm32(32, 128, eps=1e-05, affine=True)
    (1): SiLU()
    (2): Conv2d(128, 6, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
```

(3) Now we need to implement the posterior sampling of DDPM as shown in line 4 in Algorithm 2. As you can notice that in the original sampling algorithm of DDPM, the model is trained to predict the noise ϵ .

Our pretrained model takes x_t and t as input and predict the noise ϵ .

What's more, our model is also trained to predict the variance σ_t . The model output is a torch tensor of shape (B, C, H, W) , where B is the batch size, C is the number of channels and H, W are the height and width respectively. C here for our model is 6, with the first 3 channels for the noise prediction ϵ and the last 3 channels for σ_t .

Implement the `p_sample` function of `DDPMDiffusion` in `guided_diffusion/simple_diffusion.py` for unconditional posterior sampling. Follow the sampling algorithm. Attach your code to the report.

(4) Now we have everything we need to perform unconditional sampling!

Implement the `p_sample_loop` of `DDPMDiffusion` in `simple_diffusion.py` for unconditional sampling, using the DDPM sampling algorithm.

(6) Run the code below to see what we can get from unconditional distillation. Include your results in your report.

```
In [5]: from guided_diffusion.simple_diffusion import *
import torchvision
from util.img_utils import clear_color, mask_generator
from torchvision.transforms.functional import to_pil_image, pil_to_tensor
from util.img_utils import clear_color, mask_generator
from PIL import Image
```

```
In [6]: diffusion_config = {
    'sampler': 'ddpm',
    'steps': 1000,
    'noise_schedule': 'linear',
    'model_mean_type': 'epsilon',
    'model_var_type': 'learned_range',
    'dynamic_threshold': False,
    'clip_denoised': True,
    'rescale_timesteps': False,
    'timestep_respacing': 1000}

sampler = create_sampler(**diffusion_config) # Instantiate DDPMdiffusion
sample_fn = partial(sampler.p_sample_loop, model=model, measurement_cond_fn=)

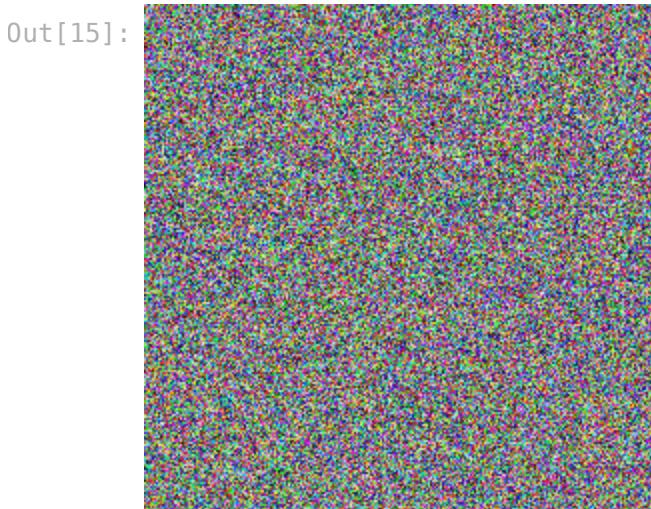
x_start = torch.randn((1, 3, 256, 256), device=device)
out_path = os.path.join(save_dir, 'uncond_ddpm')
```

```
In [13]: sample = sample_fn(x_start=x_start, measurement=None, record=True, save_root
```

0%		0/1000 [00:00<?, ?it/s]
100%		1000/1000 [00:53<00:00, 18.86it/s]

We start from random noise of the same size as our output image:

```
In [15]: torchvision.transforms.functional.to_pil_image((x_start[0] + 1)/2)
```



We can then generate human faces with our pretrained model when given the noise input:

```
In [16]: plt.imsave(os.path.join(out_path, 'output', '0.png'), clear_color(sample))
torchvision.transforms.functional.to_pil_image((sample[0] + 1)/2)
```

Out[16]:



Task 2: Unconditional Sampling with DDIM

In this task, you will implement an improved sampling algorithm from Denoising Diffusion Implicit Models(DDIM) paper. DDIM sampling applies an improved update rule to sample from $p(x_{t-1}|x_t, x_0)$. The update rule is given by

Leveraging the above improved update rule, DDIM can be used to accelerate the sampling algorithm by only using a subset of the timesteps as before.

In `simple_diffusion.py` update the method `p_sample` under the class `DDIMDiffusion`, to implement the above update rule for DDIM sampling.

```
In [16]: timestep_spacing = 100

diffusion_config = {
    'sampler': 'ddim',
    'steps': 1000,
    'noise_schedule': 'linear',
    'model_mean_type': 'epsilon',
    'model_var_type': 'learned_range',
    'dynamic_threshold': False,
    'clip_denoised': True,
    'rescale_timesteps': True,
    'timestep_respacing': f'ddim{timestep_spacing}'}

sampler = create_sampler(**diffusion_config)
sample_fn = partial(sampler.p_sample_loop, model=model, measurement_cond_fn=
```

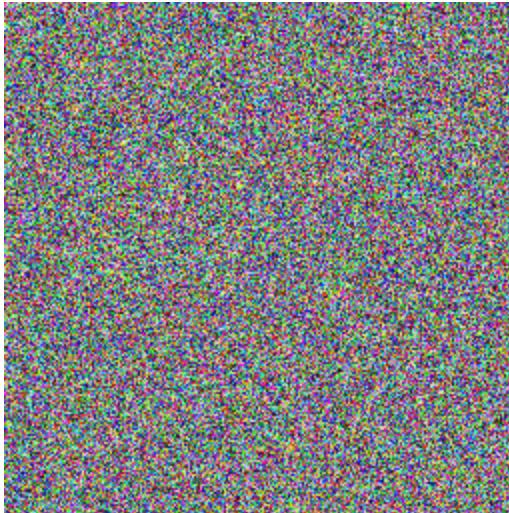
```
x_start = torch.randn((1, 3, 256, 256), device=device)
out_path = os.path.join(save_dir, 'uncond_ddim')
save_path = os.path.join(out_path, "progress")
os.makedirs(save_path, exist_ok=True)
```

In [12]: `sample = sample_fn(x_start=x_start, measurement=None, record=True, save_root`

100%|██████████| 100/100 [00:05<00:00, 17.73it/s]

In [13]: `x_start_plot = to_pil_image((x_start[0] + 1)/2)
x_start_plot`

Out[13]:



In [14]: `sample_plot = to_pil_image((sample[0] + 1)/2)
sample_plot`

Out[14]:



Task 3: Inverse problem with RePaint

In this task, you will be applying the generative DDPM to solve an interesting problem of Image Inpainting. Image inpainting refers to filling out regions of the image that are unknown apriori. Here, we assume that a mask m indicating the known region is given to us.

Repaint Diffusion applies an update rule to the input image as shown below,

where the known region is sampled using

$$x_{t-1}^{known} \sim N(\bar{\alpha}_t x_0, (1 - \bar{\alpha}_t) I)$$

and the unknown region is sampled from the diffusion model as

$$x_{t-1}^{unknown} \sim N(\mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

and the new sample x_{t-1} is obtained using

$$x_{t-1} = m \odot x_{t-1}^{known} + (1 - m) \odot x_{t-1}^{unknown}$$

In `simple_diffusion.py` update the method `p_sample` under the class `Repaint`, to implement the above update rule for inpainting.

The summarized algorithm is given by

```
In [23]: repaint_beta = get_named_beta_schedule('linear', 1000)

model_config = {
    'image_size': 256,
    'num_channels': 128,
    'num_res_blocks': 1,
    'channel_mult': '',
    'learn_sigma': True,
    'class_cond': False,
    'use_checkpoint': False,
    'attention_resolutions': 16,
    'num_heads': 4,
    'num_head_channels': 64,
    'num_heads_upsample': -1,
    'use_scale_shift_norm': True,
    'dropout': 0.0,
    'resblock_updown': True,
    'use_fp16': False,
    'use_new_attention_order': False,
    'model_path': 'models/ffhq_10m.pt'
}

diffusion_config = {
    'sampler': 'repaint',
    'steps': 1000,
    'noise_schedule': 'linear',
    'model_mean_type': 'epsilon',
    'model_var_type': 'learned_range',
    'dynamic_threshold': False,
```

```

'clip_denoised': True,
'rescale_timesteps': True,
'timestep_respacing': 250}

repaint_conf = {
    "name": "face_example",
    "inpa_inj_sched_prev": True,
    "n_jobs": 1,
    "print_estimated_vars": True,
    "inpa_inj_sched_prev_cumnoise": False,
    "class_cond": False,
    "schedule_jump_params": {
        "t_T": 250,
        "n_sample": 1,
        "jump_length": 10,
        "jump_n_sample": 10,
    }
}

repaint_model = create_model(betas=repaint_beta, **model_config)
repaint_model = repaint_model.to(device)
repaint_model.eval()
print("Model Loaded")

gt_path = "data/datasets/gts/face/000000.png"
gt_mask_path = "data/datasets/gt_keep_masks/face/000000.png"

```

pretrained model loaded!

Model Loaded

```

In [24]: model_kwargs_keys = ['gt', 'gt_keep_mask']

pil_gt_image = Image.open(gt_path)
gt_tensor = (pil_to_tensor(pil_gt_image) / 127.5 - 1.0).to(device = 'cuda')

pil_gt_mask = Image.open(gt_mask_path)
gt_mask_tensor = (pil_to_tensor(pil_gt_mask) / 255.0).to(device = 'cuda')

model_kwargs = {
    'gt': gt_tensor,
    'gt_keep_mask': gt_mask_tensor
}

```

```

In [25]: sampler = create_sampler(**diffusion_config)
sample_fn = partial(sampler.p_sample_loop, model=repaint_model, shape=(1, 3, 256, 256))

x_start = torch.randn((1, 3, 256, 256), device=device)
out_path = os.path.join(save_dir, 'repaint')
save_path = os.path.join(out_path, "progress")
os.makedirs(save_path, exist_ok=True)

```

```
In [26]: pil_gt_image
```

Out[26]:



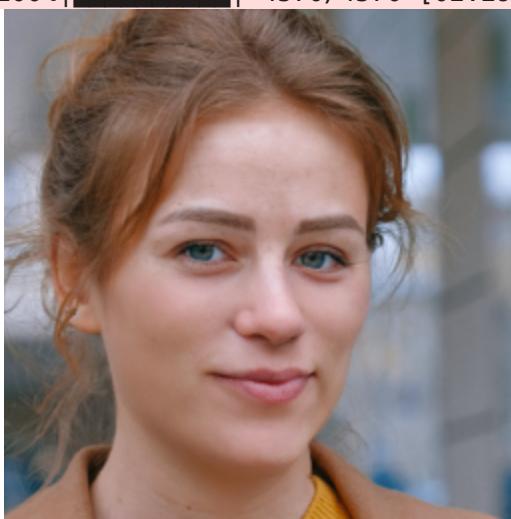
In [27]: pil_gt_mask

Out[27]:

In [28]: sample = sample_fn(noise=x_start, progress = True)
to_pil_image((sample[0] + 1)/2)

100% |██████████| 4570/4570 [02:29<00:00, 30.51it/s]

Out[28]:



Task 4: Inverse problem with DPS

In this task, you will implement the sampling algorithm with posterior sampling using a pre-trained diffusion model. This is pretty much similar to the unconditional sampling algorithm except that we now are given the corrupted input. Let's implement the algorithm below.

The algorithm follows the same process as unconditional sampling. The only difference here is that we need to use the prior provided by the diffusion model and optimize the **image** directly by taking the derivative in line 7.

- (1) We have already implemented the conditional sampling part of `p_sample_loop` and `p_sample` of `DPSDiffusion` in `guided_diffusion/simple_diffusion.py` for conditional posterior sampling. But we do encourage you to take a look at that into that.
- (2) You will need to implement the `PosteriorSampling` in `guided_diffusion/condition_methods.py`.
- (3) Run the following code. You should report one sample(including the raw image, corrupted image input and the algorithm output) of the inpainting task.
- (4) [Optional] Play around with other task configurations and operate the algorithm to see how the results look like. Report one sample(including the raw image, corrupted image input and the algorithm output) of the following task: motion deblur, gaussain deblur and super resolution. Compare the results and discuss how the algorithm perform in each task.
[Hint: Change `task_config` to play with different tasks]

```
In [29]: from guided_diffusion.measurements import get_noise, get_operator
from guided_diffusion.condition_methods import get_conditioning_method
from util.img_utils import clear_color, mask_generator
from PIL import Image
```

```
In [30]: # Prepare dataloader
# data_config = task_config['data']
data_config = {
    'name': 'ffhq',
    'root': './data/samples/'}
transform = transforms.Compose([transforms.ToTensor(),
                               transforms.Normalize((0.5, 0.5, 0.5), (0.5,
dataset = get_dataset(**data_config, transform=transform)
loader = get_dataloader(dataset, batch_size=1, num_workers=0, train=False)

# configuration of inpainting task
task_config_inpainting = {'conditioning':
                           {'method': 'ps',
                            'params': {'scale': 0.5}},
                           'measurement':
                           {'operator': {'name': 'inpainting'},
                            'mask_opt':
                            {'mask_type': 'random',
                             'mask_prob_range': (0.3, 0.7)},
```

```

        'image_size': 256},
        'noise': {'name': 'gaussian', 'sigma': 0.05}}
    }

# configuration of motion-deblur task
task_config_motion_deblur = {'conditioning':
    {'method': 'ps',
     'params': {'scale': 0.3}},
    'measurement':
    {'operator': {
        'name': 'motion_blur',
        'kernel_size': 61,
        'intensity': 0.5},
     'noise': {'name': 'gaussian', 'sigma': 0.05}}
}

# configuration of gaussian-deblur task
task_config_gaussian_deblur = {'conditioning':
    {'method': 'ps',
     'params': {'scale': 0.3}},
    'measurement':
    {'operator': {
        'name': 'gaussian_blur',
        'kernel_size': 61,
        'intensity': 3.0},
     'noise': {'name': 'gaussian', 'sigma': 0.05}}
}

# configuration of super resolution task
task_config_super_resolution = {'conditioning':
    {'method': 'ps',
     'params': {'scale': 0.3}},
    'measurement':
    {'operator': {
        'name': 'super_resolution',
        'in_shape': (1, 3, 256, 256),
        'scale_factor': 4},
     'noise': {'name': 'gaussian', 'sigma': 0.05}}
}

task_config = task_config_inpainting

measure_config = task_config['measurement']
operator = get_operator(device=device, **measure_config['operator'])
noiser = get_noise(**measure_config['noise'])
logger.info(f"Operation: {measure_config['operator']]['name'] / Noise: {measur
e_config['noise']['name']}")

# Prepare conditioning method
cond_config = task_config['conditioning']
cond_method = get_conditioning_method(cond_config['method'], operator, noise
r)
measurement_cond_fn = cond_method.conditioning
logger.info(f"Conditioning method : {task_config['conditioning']['method']}")

diffusion_config = {
    'sampler': 'dps',
    'steps': 1000,
}

```

```

'noise_schedule': 'linear',
'model_mean_type': 'epsilon',
'model_var_type': 'learned_range',
'dynamic_threshold': False,
'clip_denoised': True,
'rescale_timesteps': False,
'timestep_respacing': 1000}

sampler = create_sampler(**diffusion_config)
sample_fn = partial(sampler.p_sample_loop, model=model, measurement_cond_fn=)

out_path = os.path.join(save_dir, measure_config['operator']['name'])
os.makedirs(out_path, exist_ok=True)
for img_dir in ['input', 'recon', 'progress', 'label']:
    os.makedirs(os.path.join(out_path, img_dir), exist_ok=True)

```

2024-04-09 01:15:50,275 [DPS] >> Operation: inpainting / Noise: gaussian
2024-04-09 01:15:50,277 [DPS] >> Conditioning method : ps
DPS Initialized!

```

In [31]: if measure_config['operator']['name'] == 'inpainting':
            mask_gen = mask_generator(
                **measure_config['mask_opt'])

        for i, ref_img in enumerate(loader):
            logger.info(f"Inference for image {i}")
            fname = str(i).zfill(5) + '.png'
            ref_img = ref_img.to(device)

            # Exception) In case of inpaiting,
            if measure_config['operator']['name'] == 'inpainting':
                mask = mask_gen(ref_img)
                mask = mask[:, 0, :, :].unsqueeze(dim=0)
                measurement_cond_fn = partial(cond_method.conditioning, mask=mask)
                sample_fn = partial(sample_fn, measurement_cond_fn=measurement_cond_fn)

            # Forward measurement model (Ax + n)
            y = operator.forward(ref_img, mask=mask)
            y_n = noiser(y)

        else:
            # Forward measurement model (Ax + n)
            y = operator.forward(ref_img)
            y_n = noiser(y)

            # Sampling
            x_start = torch.randn(ref_img.shape, device=device).requires_grad_()
            sample = sample_fn(x_start=x_start, measurement=y_n, record=True, save_dir=out_path)

            plt.imsave(os.path.join(out_path, 'input', fname), clear_color(y_n))
            plt.imsave(os.path.join(out_path, 'label', fname), clear_color(ref_img))
            plt.imsave(os.path.join(out_path, 'recon', fname), clear_color(sample))

        break

```

```
2024-04-09 01:15:54,288 [DPS] >> Inference for image 0  
100%|██████████| 1000/1000 [01:50<00:00, 9.07it/s]
```

Now let's visualize some results.

Runnig the code below to visualize the raw image:

```
In [32]: Image.open(os.path.join(out_path, 'label', '00000.png'))
```

Out[32]:



And here is the corrupted image by random masks:

```
In [33]: Image.open(os.path.join(out_path, 'input', '00000.png'))
```

Out[33]:



Now let's see how our algorithm works:

```
In [34]: Image.open(os.path.join(out_path, 'recon', '00000.png'))
```

Out[34]:

